

Shiny Server Professional v1.4.0 Administrator's Guide

- 1 Getting Started
 - 1.1 Introduction
 - 1.2 System Requirements
 - 1.3 Installation
 - 1.3.1 Ubuntu (12.04+)
 - 1.3.2 RedHat/CentOS (5.4+)
 - 1.3.3 SUSE Linux Enterprise Server (11+)
 - 1.3.4 Install Shiny
 - 1.3.5 R Installation Location
 - 1.4 Stopping and Starting
 - 1.4.1 systemd (RedHat 7, Ubuntu 15.04+, SLES 12+)
 - 1.4.2 Upstart (Ubuntu 12.04 through 14.10, RedHat 6)
 - 1.4.3 init.d (RedHat 5, SLES 11)
- 2 Server Management
 - 2.1 Default Configuration
 - 2.2 Server Hierarchy
 - 2.2.1 Server
 - 2.2.2 Location
 - 2.2.3 Application
 - 2.3 run_as
 - 2.3.1 :HOME_USER:
 - 2.3.2 Running Shiny Server with Root Privileges
 - 2.3.3 :AUTH_USER:
 - 2.4 PAM Sessions
 - 2.4.1 Session Profile
 - 2.5 r_path
 - 2.6 Local App Configurations
 - 2.7 Hosting Model
 - 2.7.1 Host a Directory of Applications
 - 2.7.2 Host a Single Application
 - 2.7.3 Host Per-User Application Directories
 - 2.8 Redirecting
 - 2.9 Virtual Hosts
 - 2.10 Custom Templates
 - 2.11 Server Log
 - 2.11.1 Access Logs
 - 2.12 Environment Variable Settings
 - 2.12.1 SHINY_LOG_LEVEL
 - 2.12.2 R
 - 2.12.3 SHINY_DATA_DIR
- 3 Deploying Applications
 - 3.1 Schedulers & Application Restarts
 - 3.1.1 Restarting an Application
 - 3.1.2 Simple Scheduler
 - 3.1.3 Utilization Scheduler
 - 3.2 R Markdown
 - 3.3 Application Timeouts
 - 3.4 Session Timeouts
 - 3.5 Logging and Analytics
 - 3.5.1 Application Error Logs
 - 3.5.2 Google Analytics
 - 3.6 Program Supervisors
 - 3.7 Memory Limits
 - 3.8 Reactivity Log
 - 3.9 Specifying Protocols
 - 3.9.1 Disabling WebSockets on the Server
- 4 Authentication & Security
 - 4.1 Authentication Overview
 - 4.2 Auth Duration
 - 4.3 Flat-File Authentication
 - 4.3.1 sspasswd
 - 4.4 Google Authentication
 - 4.4.1 Create a Google Application
 - 4.4.2 Configure Shiny Server

- 4.4.3 Securing the Client Secret
 - 4.4.4 Email Address Restrictions
- 4.5 PAM Authentication
 - 4.5.1 PAM Basics
 - 4.5.2 Default PAM Configuration
 - 4.5.3 PAM and Kerberos
- 4.6 LDAP and Active Directory
 - 4.6.1 base_bind
 - 4.6.2 Securing the LDAP password
 - 4.6.3 user_bind_template
 - 4.6.4 group_search_base
 - 4.6.5 group_filter
 - 4.6.6 group_name_attribute
 - 4.6.7 trusted_ca
 - 4.6.8 check_ssl_ca
 - 4.6.9 user_filter
 - 4.6.10 user_search_base
- 4.7 Required Users & Groups
 - 4.7.1 User Authentication
 - 4.7.2 Group Authentication
- 4.8 SSL
- 4.9 Proxied Headers
- 5 Monitoring the Server
 - 5.1 Admin
 - 5.1.1 Configuration
 - 5.1.2 Organization
 - 5.1.3 Killing Processes and Connections
 - 5.2 Graphite
 - 5.3 Health Check Endpoint
 - 5.3.1 Customizing Responses
 - 5.3.2 Providing Multiple Health-Check Endpoints
- 6 Licensing & Activation
 - 6.1 Product Activation
 - 6.1.1 Activation Basics
 - 6.2 Connectivity Requirements
 - 6.2.1 Proxy Servers
 - 6.2.2 Offline Activation
- 7 Appendix
 - 7.1 Quick Start
 - 7.1.1 Host a directory of applications
 - 7.1.2 Let users manage their own applications
 - 7.1.3 Run Shiny Server on multiple ports
 - 7.1.4 Require user authentication on an application
 - 7.1.5 Host a secure Shiny Server
 - 7.1.6 Host an Application Supported by Multiple R Processes
 - 7.2 Configuration Settings
 - 7.3 Migrating & Updating
 - 7.3.1 Upgrading from Shiny Server 0.3.x
 - 7.3.2 Upgrading from Shiny Server 0.4.x
 - 7.4 Frequently Asked Questions
 - 7.4.1 What are the hardware requirements for running Shiny Server?
 - 7.4.2 Does Shiny Server work on Amazon Elastic Compute Cloud (EC2)?
 - 7.4.3 Can I Deploy Shiny Server on an Isolated Network?
 - 7.4.4 How Do I Translate My 'application' Settings to Nested 'location's?
 - 7.4.5 When I open a lot of Shiny Docs at once in Firefox, why do some not open?
 - 7.4.6 Can I Bulk-Import in Flat-File Authentication?
 - 7.4.7 How Can I Set The trusted_ca For My LDAP Server Over SSL?

1 Getting Started

1.1 Introduction

Shiny Server enables users to host and manage Shiny applications on the Internet. Shiny is an R package which uses a reactive programming model to simplify the development of R-powered web applications. Shiny Server can manage R processes running various Shiny applications over different URLs and ports. Using Shiny Server offers a variety of benefits over simply running

Shiny in R directly. These features allow the administrator to:

- Host multiple applications simultaneously, each at its own URL.
- Support web browsers which don't support WebSocket, including Internet Explorer 8 & 9.
- Enable system users to develop and manage their own Shiny applications.
- Automatically have R processes that crash or are terminated restart for the next user requesting the application.

PRO

This manual describes Shiny Server Professional which offers, among other things, the following additional features:

- Ensure your applications are protected and can only be accessed by **specific, authenticated users**.
- Scale a Shiny application to support many users by empowering a Shiny application to be backed by multiple R Shiny processes simultaneously.
- Gain insight into the performance and usage of your Shiny applications by monitoring them using a web dashboard.
- Securely encrypt data being sent to and from your applications using SSL.
- Understand and manage current and historical application resource utilization to better configure and optimize your applications.
- Fine-tune the resources devoted to each user of an application by configuring multi-process Shiny applications based on the number of concurrent sessions.
- Monitor the health of your Shiny Server using the health check endpoint.

1.2 System Requirements

Shiny Server is currently only supported on the Linux operating system. We officially support the following distributions:

- RedHat Enterprise Linux (and CentOS) 5.4+
- RedHat Enterprise Linux (and CentOS) 6.0+
- RedHat Enterprise Linux (and CentOS) 7.0+
- Ubuntu 12.04
- Ubuntu 13.04
- Ubuntu 14.04
- Ubuntu 15.04
- SUSE Linux Enterprise 11 SP1 and higher

We currently only support the x86-64 architecture. As described in the [Installation](#) section, you will **install R and the Shiny** package prior to **installing Shiny Server**. **Root** privileges will be required both in the installation process and also at runtime.

PRO

Shiny Server Professional allows you to host multiple R processes concurrently to balance the load of incoming requests across multiple Shiny instances. The [Utilization Scheduler](#) section explains how to configure this in more detail. We recommend carefully evaluating the computational properties and memory profile of a Shiny application using the tools discussed in the [Admin](#) section. This dashboard will provide much more insight into the profile of a Shiny application than was previously possible.

1.3 Installation

The Shiny Server installer does not come with R, nor does it provide the Shiny R package. Below are the steps for installing each of them separately.

If you had previously installed Shiny Server 0.3.x or 0.4.x, please see the relevant section before proceeding to ensure a seamless upgrade:

- [Upgrading from Shiny Server 0.3.x](#)
- [Upgrading from Shiny Server 0.4.x](#)

1.3.1 Ubuntu (12.04+)

Installing R

Shiny Server recommends an installation of R version 3.0 or higher. To install the latest version of R you should first add the CRAN repository to your system as described here:

- Debian: <http://cran.rstudio.com/bin/linux/debian/README.html>
- Ubuntu: <http://cran.rstudio.com/bin/linux/ubuntu/README.html>

You can then install R using the following command:

```
$ sudo apt-get install r-base
```

NOTE: if you do not add the CRAN Debian or Ubuntu repository as described above this command will install the version of R corresponding to your current system version. Since this version of R may be a year or two old it is strongly recommended that you add the [CRAN repositories](#) so you can run the most up to date version of R.

At this point, follow the instructions in [Install Shiny](#) to setup the necessary packages in R.

Once Shiny is installed, you can begin the installation of Shiny Server. You should have been provided with a .deb installer for Shiny Server. If you only have a link to this file, you can use `wget` to download the file to the current directory. You will need to first install the `gdebi-core` package to install Shiny Server and its dependencies. Once the .deb file is available locally, run the following commands to complete the installation of Shiny Server.

```
sudo apt-get install gdebi-core
sudo gdebi shiny-server-1.4.0.deb
```

This will install Shiny Server into `/opt/shiny-server/`, with the main executable in `/opt/shiny-server/bin/shiny-server` and also create a new shiny user. The rest of this guide will document the intricacies of how you can configure and manage Shiny Server. If you're just looking to get up and running quickly, we recommend you look at the [Quick Start](#) section in the appendix which will walk you through the process of installing and configuring a Shiny application.

1.3.2 RedHat/CentOS (5.4+)

Prerequisites

Shiny Server recommends an installation of R version 3.0 or higher. Shiny Server has several dependencies on packages (including R itself) found in the Extra Packages for Enterprise Linux (EPEL) repository. If you don't already have this repository available you should add it to your system using the instructions found here: <https://fedoraproject.org/wiki/EPEL>. On some distributions of RedHat/CentOS, the R package references dependencies that are not available by default. In this case, you may need to edit the `/etc/yum.repos.d/redhat.repo` file to enable the `rhel-6-server-optional-rpms` (by setting `enabled = 1`) before you can install the R package.

After enabling EPEL you should then ensure that you have installed the version of R available from EPEL. You can do this using the following command:

```
$ sudo yum install R
```

At this point, follow the instructions in [Install Shiny](#) to setup the necessary packages in R.

Once Shiny has been installed, you can begin the installation of Shiny Server. You should have been provided with an RPM file which contains Shiny Server and all of its dependencies (other than R and Shiny). You can install this rpm file using `yum`. If you have only a link to the RPM file, you can use `wget` to download the file to the current directory. You can install this RPM file using `yum`.

```
$ sudo yum install --nogpgcheck shiny-server-1.4.0.rpm
```

This will install Shiny Server into `/opt/shiny-server/`, with the main executable in `/opt/shiny-server/bin/shiny-server` and also create a new shiny user. The rest of this guide will document the intricacies of how you can configure and manage Shiny Server. If you're just looking to get up and running quickly, we recommend you look at the [Quick Start](#) section in the appendix which will walk you through the process of installing and configuring a Shiny application.

1.3.3 SUSE Linux Enterprise Server (11+)

Prerequisites

Shiny Server recommends an installation of R version 3.0 or higher. You can install R for SUSE Linux Enterprise Server using the instructions on the CRAN SUSE download page: <http://cran.rstudio.com/bin/linux/suse/>. Note that the binaries suggested here have one additional requirement that isn't satisfied using the default repositories. Before installing R you should install the `libgfortran43` package. This package is available from the [SUSE Linux Enterprise SDK](#). If the SDK repository is available in your environment you can install `libgfortran43` as follows:

```
$ sudo zypper install libgfortran43
```

At this point, follow the instructions in [Install Shiny](#) to setup the necessary packages in R.

Once R and Shiny have been installed, you can begin the installation of Shiny Server.

Additional Prerequisites for SLES

The Shiny Server Professional RPM binaries link to version 0.9.8 of the openssl and libcrypto shared libraries. On RedHat systems the shared library files have a suffix of ".6" (e.g. libcrypto.so.6) whereas on SLES they have a suffix of "0.9.8". As a result you should soft-link the existing 0.9.8 shared libraries to a version with a .6 suffix by running the following commands:

```
$ sudo ln -f -s /usr/lib64/libssl.so.0.9.8 /usr/lib64/libssl.so.6
$ sudo ln -f -s /usr/lib64/libcrypto.so.0.9.8 /usr/lib64/libcrypto.so.6
```

The Shiny Server Professional metrics and monitoring features use a set of Python scripts to introspect the system for performance information. As part of installation a Python native extension needs to be built, and this requires the python-devel-26 package available from the [SUSE Linux Enterprise SDK](#). If the SDK repository is available in your environment you can install python-devel-26 as follows:

```
$ sudo zypper install python-devel-26
```

Installation

You should have been provided with an RPM file which contains Shiny Server and all of its dependencies (other than R and Shiny). Note that the CentOS 5 RPM of Shiny Server is the one that's compatible with SLES, so know that there's not a problem if you see that branding on your file. You can install this RPM file using zypper. If you have only a link to the RPM file, you can use wget to download the file to the current directory. You can install this RPM file using zypper.

```
$ sudo zypper install shiny-server-1.4.0.rpm
```

During the installation zypper will prompt to ensure that you still want to install even without the libffi dependency. The prompt will look like this:

```
Problem: nothing provides libffi needed by shiny-server
Solution 1: do not install shiny-server
Solution 2: break shiny-server by ignoring some of its dependencies
Choose from above solutions by number or cancel [1/2/c] (c): 2
```

The libffi dependency is not required on SLES so it is safe to install without it. You should enter "2" (ignore the dependencies) at this prompt and continue with the installation.

This will install Shiny Server into /opt/shiny-server/, with the main executable in /opt/shiny-server/bin/shiny-server and also create a new shiny user. The rest of this guide will document the intricacies of how you can configure and manage Shiny Server. If you're just looking to get up and running quickly, we recommend you look at the [Quick Start](#) section in the appendix which will walk you through the process of installing and configuring a Shiny application.

1.3.4 Install Shiny

Before Shiny Server can be installed, the Shiny package must be installed in the system library; you typically need sudo privileges to install to this library.

Setup the Server for Secure Package Installation

R packages can be installed securely using an HTTPS CRAN mirror, or insecurely using an HTTP mirror. Beginning with R version 3.2.2, HTTPS is the default preference when installing packages, but older versions of R default to insecure HTTP. You can change this behavior in older versions of R by setting the download.file.method option in your .Rprofile. See [R Startup Files](#) for details on where these files are located.

You should add the following line to Rprofile.site to affect how all users install packages, or the ~/.Rprofile file for individual users who will be running Shiny applications in Shiny Server.

For R 3.2

```
options(download.file.method = "libcurl")
```

For R 3.1 and earlier

```
options(download.file.method = "wget")
```

You could alternatively use curl instead of wget, if wget is not available on your server.

You'll also want to specify a secure default CRAN mirror in this same file. You can do that using the following code:

```
local({  
  r <- getOption("repos")  
  r["CRAN"] <- "https://cran.rstudio.com/"  
  options(repos=r)  
})
```

Install The Shiny Package

Shiny Server currently requires Shiny version 0.7.0 or later. The following command will download and install the shiny package from CRAN in the system library.

```
$ sudo su - -c "R -e \"install.packages('shiny')\""
```

If you have not set a default CRAN repository (as described in the section above) you may need to specify the repos parameter from which to download shiny and its dependencies.

Once this command completes, you can continue with the installation of Shiny Server.

1.3.5 R Installation Location

Shiny Server expects that **R is available as** an executable named R and is in the PATH of the user which you run shiny-server as. Note that on some CentOS systems, the PATH will be overridden by the startup script to /sbin:/usr/sbin:/bin:/usr/bin. On such systems, if R is not available in one of these locations (regardless of the user's PATH), you'll need to adjust the startup script.

To allow Shiny Server to search for R in additional locations, you'll alter the file in /etc/init.d/shiny-server or /etc/init/shiny-server.conf depending on which startup system you're using (as discussed in [Stopping and Starting](#)). You can either adjust the PATH variable to include the directory where R will be found, or you can set an environment variable named R to tell Shiny Server exactly where it should look for the executable.

If you choose to adjust the PATH, you can add the directory in which the executable named R is found to the line that defines the PATH environment variable (PATH=/sbin:/usr/sbin:/bin:/usr/bin).

If instead you choose to tell Shiny Server the exact executable to run (which is necessary if the executable is not named R on your system), you'll need to define a new environment variable named R. You can do this by adding a line that looks something like env R=/usr/local/bin/R-3-0-1 for Upstart, or export R=/usr/local/bin/R-3-0-1 for init.d.

1.4 Stopping and Starting

The installer will automatically deploy the necessary scripts to ensure that Shiny Server is started automatically on boot. When possible, we use the Upstart system to manage the shiny-server service. If Upstart is not available, we will deploy an init.d script to automatically start and stop the service.

1.4.1 systemd (RedHat 7, Ubuntu 15.04+, SLES 12+)

systemd is a management and configuration platform for Linux. The newest versions of most major Linux distributions have adopted systemd as their default init system.

The Shiny Server installer will automatically install a systemd service called shiny-server, which will cause the shiny-server program to be started and stopped automatically when the machine boots up and shuts down. The shiny-server service will also be launched automatically when the installer has successfully installed the program.

To manually start or stop the server, you can use the following commands.

```
$ sudo systemctl start shiny-server
```

```
$ sudo systemctl stop shiny-server
```

You can restart the server with:

```
$ sudo systemctl restart shiny-server
```

This command will shutdown all running Shiny processes, disconnect all open connections, and re-initialize the server.

If you wish to keep the server and all Shiny processes running without interruption, but reload the configuration, you can use the `systemctl` command to send a `SIGHUP` signal:

```
$ sudo systemctl kill -s HUP --kill-who=main shiny-server
```

This will cause the server to re-initialize but will not interrupt the current processes or any of the open connections to the server.

You can check the status of the shiny-server service using:

```
$ sudo systemctl status shiny-server
```

And finally, you can use the `enable/disable` commands to control whether Shiny Server should be run automatically at boot time:

```
$ sudo systemctl enable shiny-server
```

```
$ sudo systemctl disable shiny-server
```

1.4.2 Upstart (Ubuntu 12.04 through 14.10, RedHat 6)

Upstart is a system used to automatically start, stop and manage services. The installer will automatically deploy an Upstart script to `/etc/init/shiny-server.conf`. This script will initialize shiny-server as soon as the network is activated on the machine and stop when the machine is being shut down.

The Upstart script will also ensure that shiny-server is respawned if the process is terminated unexpectedly. However, in the event that there is an issue which will consistently prevent Shiny Server from being able to start (such as a bad configuration file), Upstart will give up on restarting the service after approximately 5 failed attempts within a few seconds. For this reason, you may see multiple repetitions of a bad Shiny Server startup attempt before it transitions to the stopped state.

To start or stop the server, run the following commands, respectively.

```
$ sudo start shiny-server
```

```
$ sudo stop shiny-server
```

To restart the server you can run:

```
$ sudo restart shiny-server
```

This command will shutdown all running Shiny processes, disconnect all open connections, and re-initialize the server. Note that restart *will not* reread the Upstart definition at `/etc/init/shiny-server.conf`. So if you have changed, for instance, some environment variables in that file, you will need to stop and start to have those changes take effect.

If you wish to **keep the server and all Shiny processes running without interruption, but reload the configuration**, you can use the `reload` command as in:

```
$ sudo reload shiny-server
```

This will cause the server to re-initialize but will not interrupt the current processes or any of the open connections to the server.

Known Bug: Due to a bug in the version of Upstart which comes with Ubuntu 13.04, reload will not behave as expected on that platform and should not be used.

To check the status or retrieve the Process ID associated with shiny-server, run the following:

```
$ status shiny-server
```

1.4.3 init.d (RedHat 5, SLES 11)

On some older Operating Systems (such as RedHat 5) the Upstart system may not be available. These systems will require the use of a script in /etc/init.d/ to start and manage the shiny-server daemon. This script does not have the capability to automatically restart the shiny-server daemon if it terminates unexpectedly. The installer uses chkconfig to ensure that shiny-server will start automatically on boot.

To start or stop the server, use the following commands, respectively:

```
$ sudo /sbin/service shiny-server start
```

```
$ sudo /sbin/service shiny-server stop
```

To restart the server, you can run:

```
$ sudo /sbin/service shiny-server restart
```

If you wish to keep the server and all Shiny processes running without interruption, but reload the configuration, you can use reload:

```
$ sudo /sbin/service shiny-server reload
```

This will cause the server to re-read the configuration file but will not interrupt the current processes and outstanding connections to the server.

To check the status or retrieve the Process ID associated with shiny-server, run the following:

```
$ /sbin/service shiny-server status
```

2 Server Management

2.1 Default Configuration

Initially, Shiny Server uses the following configuration file. Some users will find that this configuration meets their needs; others may find it useful to create a custom configuration. Details about each available setting and parameter are available in the [Appendix](#). As a brief introduction, however, the following configuration file will be installed at `/etc/shiny-server/shiny-server.conf` if it doesn't already exist:

```
# Define the user we should use when spawning R Shiny processes
run_as shiny;

# Define a top-level server which will listen on a port
server {
  # Instruct this server to listen on port 3838
  listen 3838;

  # Define the location available at the base URL
  location / {
    #### PRO ONLY ####
    # Only up to 20 connections per Shiny process and at most 3 Shiny processes
    # per application. Proactively spawn a new process when our processes reach
```



```

# 90% capacity.
utilization_scheduler 20 .9 3;
#### END PRO ONLY ####

# Run this location in 'site_dir' mode, which hosts the entire directory
# tree at '/srv/shiny-server'
site_dir /srv/shiny-server;

# Define where we should put the log files for this location
log_dir /var/log/shiny-server;

# Should we list the contents of a (non-Shiny-App) directory when the user
# visits the corresponding URL?
directory_index on;
}
}

# Setup a flat-file authentication system. {.pro}
auth_passwd_file /etc/shiny-server/passwd;

# Define a default admin interface to be run on port 4151. {.pro}
admin 4151 {
  # Only permit the user named `admin` to access the admin interface.
  required_user admin;
}

```

Lines beginning with a # are treated as comments and not parsed when configuring the server. Shiny Server can be configured to **host multiple servers on different ports or hostnames**. Each server can have locations which are capable of serving Shiny Applications and potentially static assets, as well. Individual applications can also **override the settings** applied to their parent location. These concepts are explained in further detail in the [Server Hierarchy](#) section. The default configuration above will create a single server listening on port 3838, serving any application contained within /srv/shiny-server/ at the root URL (/). Each possible setting in the configuration file is explained in the [Appendix](#).

Most users will want to customize the configuration to meet their needs. The server will load its configuration from a file stored at **/etc/shiny-server/shiny-server.conf**; it is in this location that you should customize your Shiny Server configuration.

PRO

This configuration will also create an administrative dashboard running on port **4151**. The admin interface requires that a user authenticate themselves, as discussed further in the chapter on [Authentication & Security](#). The configuration will attempt to use a flat-file authentication system stored at **/etc/shiny-server/passwd**; an empty database is created for you here during installation. To create a new user named admin in this file to allow you to login to the dashboard, execute the following command

```
$ sudo /opt/shiny-server/bin/sspasswd /etc/shiny-server/passwd admin
```

and then enter and verify the password you wish to use for this user.

2.2 Server Hierarchy

Detailed descriptions of all available parameters are available in the appendix, but it is important to understand the overall hierarchy of the Shiny Server configuration file when editing the file.

2.2.1 Server

The server setting defines an HTTP server which will listen on a port/IP address combination. For example, the following lines:

```

server {
  listen 80;
}

```

define a server that would listen on port 80. A server can also define a **server_name** to limit the virtual hostnames on which it

listens, as described in the [Virtual Hosts](#) section.

2.2.2 Location

The location setting is defined within a server and defines how a particular URL path should be served. For instance, the following settings:

```
server {
  ...
  # Define the location '/specialApp'
  location /specialApp {
    # Run this location in 'app_dir' mode, which will host a single Shiny
    # Application available at '/srv/shiny-server/myApp'
    app_dir /srv/shiny-server/myApp
  }

  # Define the location '/otherApps'
  location /otherApps {
    # Run this location in 'site_dir' mode, which hosts the entire directory
    # tree at '/srv/shiny-server/apps'
    site_dir /srv/shiny-server/apps;
  }
  ...
}
```

would define two locations, one which serves (potentially) a multitude of applications at the URL `/otherApps/`, and another which serves a single application at the URL `/specialApp/`. The various hosting models which can be applied to a location are described in the section on [Hosting Models](#).

location directives can also be nested to provide more granular settings for a particular sub-location. For instance, if you used Google Analytics but the Finance department wanted to use its own Google Analytics ID for their finance directory, you could accomplish that with a configuration like:

```
server {
  ...
  # Define the '/depts' location
  location /depts {
    # Host a directory of applications
    site_dir /srv/departmentApps;

    # Provide a default/global GAID
    google_analytics_id "UA-12345-1";

    # Define the '/finance' location.
    # Corresponds to the URL '/depts/finance', and the filesystem path
    # '/srv/departmentApps/finance' as defined by the parent location.
    location /finance {
      # Provide a custom GAID for only this sub-location
      google_analytics_id "UA-54321-9";
    }
  }
  ...
}
```

Note that the nested location's path is relative to the parent location's. Any directive that can be used for a location can be used in a nested location and will override the value specified in the parent, if any. If not overridden, the settings will be inherited by sub-locations from their parent locations.

2.2.3 Application

The application setting has been deprecated and removed from Shiny Server as of version 0.4.2.

See [How Do I Translate My 'application' Settings to Nested 'location's?](#) for help migrating.

2.3 run_as

Understanding **which user will execute the R Shiny processes** is important for a variety of reasons. For one, the paths in which R will look for packages (`.libPaths()`) are often **user-dependent**. In order to ensure that the libraries required to run a Shiny application are installed, you must first understand which user will be running the application. Additionally, directories will likely have restrictions regarding which users are able to read or write in them. The server should be configured in such a way that the user running the Shiny applications has the minimal privileges to do the work required by those applications.

locations configured with `user_apps` will be executed as the user in whose home directory the application was found. For locations configured with `site_dir`, `user_dirs`, and `app_dir`, the `run_as` setting will be used to **determine which user should spawn** the R Shiny processes. This setting can be configured globally, or for a particular server or location. For example, the following configuration:

```
location / {  
  run_as tim;  
}
```

would execute all applications contained within this scope as the user `tim` with all of `tim`'s relevant `.libPaths`.

Additionally, **R processes are spawned using the Bash login shell**. This means that prior to the execution of the R session the Bash shell will read and execute commands from this file if it exists:

```
/etc/profile
```

After reading that file, it looks for the following files and reads and executes commands from the first one that exists and is readable (it's important to note that only one of these files will be read and executed):

```
~/.bash_profile  
~/.bash_login  
~/.profile
```

So in the above configuration example, any environment variables defined in `/etc/profile` or `/home/tim/.bash_profile` (assuming `tim`'s home directory was in the standard location) would be available to the R process, since Shiny Server is configured to run applications as the user named `tim`.

2.3.1 :HOME_USER:

`:HOME_USER:` is a special instance of a `run_as` user. When combined with the `user_dirs` hosting model (described in the section entitled [Host Per-User Application Directories](#)), **this setting will instruct Shiny Server to run the process as the user in whose home directory the application exists**. For instance, an application stored in `/home/jim/ShinyApps/app1` would run as `jim`, whereas an app stored in `/home/kelly/ShinyApps/app1` would run as `kelly`.

Because this username is "special", it may make sense to provide an additional username as a "fallback" user. For instance, in the following configuration:

```
run_as :HOME_USER: shiny;  
  
...  
  
location /users {  
  user_dirs;  
}  
  
location /apps {  
  site_dir /srv/shiny-server;  
  log_dir /var/log/shiny-server;  
  directory_index on;  
}
```

all of the `user_dirs` applications hosted in `/users` would run as `:HOME_USER:` (the user in whose home directory the application exists), since that's the highest-priority `run_as` user. Because `:HOME_USER:` is only meaningful to the `user_dirs` hosting model, and not to `site_dir`, the applications hosted in `/apps` will not run as `:HOME_USER:` but as whichever user comes next in the list -- in this

case, as shiny.

Of course, you could also accomplish this same effect with the following arrangement:

```
location /users {
    run_as :HOME_USER;;
    user_dirs;
}

location /apps {
    run_as shiny;
    site_dir /srv/shiny-server;
    log_dir /var/log/shiny-server;
    directory_index on;
}
```

As described in [Host Per-User Application Directories](#), `user_apps` and `user_dirs` differ only in how they behave with respect to the `run_as` setting. `user_apps` will *always* host the apps as if `run_as` were set to `:HOME_USER::`, regardless of how `run_as` is actually set. `user_dirs`, on the other hand, will respect the `run_as` configuration.

2.3.2 Running Shiny Server with Root Privileges

Aside from spawning R Shiny processes as particular users, the shiny-server process itself can be configured to run as different users to control its privileges. There are multiple scenarios in which Shiny Server would need to be run as root:

1. If `user_apps` or `user_dirs` is enabled for any location. In order to host applications as various users, Shiny Server must have root privileges.
2. If your configuration uses `run_as` to spawn applications as multiple different users.
3. If you're running any server on a privileged port (a port in the range of 1-1024).
4. If you're using `auth_pam` to leverage PAM authentication. (Pro only)
5. If you're using `auth_google` and store your client secret in a file. (Pro only)

By default, the shiny-server process will be started as the root user, then will spawn R processes according to the corresponding `run_as` setting. You can, however, run the shiny-server process as a non-privileged user such as the shiny user if none of the limitations above are violated.

Shiny Server Professional leverages PAM to spawn sessions for users. See the section on [PAM Sessions](#) to find more details about how you can use PAM to tailor constraints on the Shiny processes that Shiny Server spawns.

For Shiny Server Pro, if it detects that it is running as root but has no need to (i.e. none of the criteria above were met), then it will drop its privileges and run as an unprivileged user; note that, for security reasons, that change is irreversible. This means that once Shiny Server drops its privileges, it wouldn't be able to regain them should a new configuration be loaded that requires such privileges.

By default, if shiny-server detects that it doesn't need to run with root privileges, it will attempt to run as an unprivileged user to be more secure. It will determine which user to run as by inspecting the `run_as` directives in your configuration file -- if there is only one user which we will be running Shiny as (and the aforementioned constraints are satisfied), the entire Shiny Server process will just run as that user. The default configuration sets `run_as` to shiny, so the process will run as the shiny user. While the process will run as the shiny user (or whatever other user you're running as), it may not behave in exactly the same way as interactively logging in as that user. In particular, the supplemental groups will not be set for the user.

Be aware that upon installing Shiny Server, permissions on the file system are set to enable the shiny user to read and write from various directories as is necessary to allow Shiny Server to run as the shiny user. Running Shiny Server as another user will require that you adjust the permissions to grant this other user the necessary privileges to run Shiny Server. In particular, ensure the user has write privileges on these paths (recursively):

- `/var/lib/shiny-server/` (or whatever custom `SHINY_DATA_DIR` setting you're using)
- `/var/log/shiny-server/` (and/or whatever other directories you use for logging)

and read privileges on these paths (recursively):

- `/srv/shiny-server/` (and/or whatever other directories you're using to host Shiny applications)
- `/opt/shiny-server/`
- `/etc/shiny-server/` (Note that you should enable only read access on this directory, as you likely don't want to allow your Shiny applications (which also run as shiny) to be able to write to your configuration or password file.)

Finally, if the directory `/tmp/shiny-server/` exists, it (and all files within it) should be owned by the user you specify.

2.3.3 :AUTH_USER: - Pro Only!

:AUTH_USER: is an additional special case of run_as available in Shiny Server Pro. :AUTH_USER: instructs Shiny Server to run applications as whatever user the visitor to the application is logged in as. For this reason, it only works when using [PAM Authentication](#) -- as that is the only authentication scheme that can ensure that the username the visitor is logged in as is a real user on the server which would be capable of running Shiny applications.

If the user is not logged in, this setting will have no effect and the next user in the run_as chain will be used. If there are no additional users, then the user will be prompted to login.

Much like :HOME_USER:, :AUTH_USER: can be part of a list of users provided to run_as. Again, the first username provided will be used if it's available; if it's not, the second will be used, and so on.

So a configuration such as the following:

```
run_as :HOME_USER: :AUTH_USER: shiny;
```

would indicate that the following priorities should be used:

1. If in the user_dirs hosting mode (the only mode which respects :HOME_USER:), then :HOME_USER: should be used to determine which user to run the applications as.
2. Otherwise, :AUTH_USER: should be used to determine who to run the applications as. Thus, if the visitor is logged in as tim, then the application will be started as the user tim. If the user is not logged in, then we'll proceed to the next option.
3. Finally, if neither of the other two could be used (i.e. we're not on the user_dirs hosting mode and the visitor is not logged in), then we'll just run the application as the shiny user.

Alternatively, in a configuration like the following:

```
run_as :AUTH_USER;;
```

We provide no "regular" user to fall back on, thus if the visitor is not logged in, no process will be started and the visitor will be prompted to login. Once logged in, we'll start a process based on whom the visitor logged in as.

Note that only one "non-special" user should be provided in the run_as configuration. It would be incorrect to use a configuration such as:

```
run_as user1 user2;
```

as both of these users are "regular" users. In this configuration, user2 would never be used, as the program would emit an error if user1 was not a valid user on the system.

There is a known issue at the time of writing in which combining :AUTH_USER: with a Utilization Scheduler will create a new scheduler for each user. This means that in the following configuration:

```
run_as :AUTH_USER:
utilization_scheduler 3 .5 2;
```

each user will be permitted to spawn two processes.

2.4 PAM Sessions - Pro Only!

Shiny Server Professional uses PAM (Pluggable Authentication Modules) for both user authentication as well to establish the resources available for R sessions. Binding resources (and limits on their use) to Shiny sessions is accomplished by calling the PAM session API. This section explains how to configure and customize PAM sessions with Shiny Server.

2.4.1 Session Profile - Pro Only!

Shiny Server supports [PAM Authentication](#), in addition to PAM sessions. PAM Authentication is described in the linked section and is used to determine the constraints around when users should be allowed to login. PAM sessions, however, are used to set the constraints and limitations on the Shiny process that a user may spawn.

The PAM directive that enables authentication with root privilege only (auth sufficient pam_rootok.so) needs to be present in the PAM profile. The behavior that Shiny Server requires is essentially same as that of the su command (impersonation of a user

without a password). Therefore by default Shiny Server uses the `/etc/pam.d/su` profile for running Shiny sessions.

2.4.1.1 Creating a Custom Profile - Pro Only!

The `/etc/pam.d/su` profile has different default behavior depending upon your version of Linux and local configuration. Depending upon what type of behavior you want associated with Shiny sessions (e.g. mounting of disks, setting of environment variables, enforcing of resource limits, etc.) you'll likely want to create a custom profile for Shiny sessions. For example, if you wanted to use a profile named `shiny-session` you would add this to the configuration file:

```
pam_sessions_profile shiny-session;
```

Here is in turn what the custom profile might contain in order to enable a few common features of PAM sessions (this is based on a modified version of the default `su` profile on Ubuntu):

`/etc/pam.d/shiny-session`

```
# This allows root to su without passwords (this is required) {.pro}
auth      sufficient pam_rootok.so

# This module parses environment configuration file(s) {.pro}
# and also allows you to use an extended config {.pro}
# file /etc/security/pam_env.conf. {.pro}
# parsing /etc/environment needs "readenv=1" {.pro}
session   required  pam_env.so readenv=1

# Locale variables are also kept into /etc/default/locale in etch {.pro}
# reading this file *in addition to /etc/environment* does not hurt {.pro}
session   required  pam_env.so readenv=1 envfile=/etc/default/locale

# Enforces user limits defined in /etc/security/limits.conf {.pro}
session   required  pam_limits.so

# The standard Unix authentication modules {.pro}
@include common-auth
@include common-account
@include common-session
```

The above serves as a good default session profile for Shiny Server. If you want to learn more about PAM profile configuration the following are good resources:

- http://www.linux-pam.org/Linux-PAM-html/Linux-PAM_SAG.html
- <http://linux.die.net/man/8/pam.d>
- <http://www.linuxjournal.com/article/2120>
- <http://www.informit.com/articles/article.aspx?p=20968>

2.5 r_path - Pro Only!

The `r_path` setting can be used to specify which version of R should be used in a particular context. Often, Shiny applications may be written with dependencies that require a particular version of R. In such a case, multiple versions of R could be installed on a server and `r_path` could be used to distinguish which applications should use which version of R.

The following example demonstrates a potential use of this setting.

```
run_as shiny;

# In general, use the R 3.0 binary provided below. {.pro}
r_path "/opt/R-3.0/bin/R";

server {
  listen 3838;
```

```
location / {
  site_dir /srv/shiny-server;
  log_dir /var/log/shiny-server;

  location /oldApp {
    # For this particular location, use an older version of R.
    r_path "/opt/R-2.14/bin/R";
  }
}
}
```

The value of `r_path` should specify the absolute path to the R executable.

2.6 Local App Configurations

In the global configuration file, the `allow_app_override` setting can be specified to enable local app configurations. If present, this setting enables owners of Shiny applications to customize the properties of their own applications using a file named `.shiny_app.conf`. This file can be placed within an application directory (alongside the `server.R` and `ui.R` files) and can contain settings which configure how Shiny Server should manage the application.

This behavior is controlled by the `allow_app_override` setting which is disabled by default. In order to enable local application configurations, you can add to your configuration file `allow_app_override true`; (or just `allow_app_override`; for short).

Server administrators should be mindful of the fact that allowing application owners to specify their own scheduler parameters has potential performance implications, as an application owner could enable his or her application to consume more resources than is desired. It is thus recommended that write privileges on the `.shiny_app.conf` files be granted judiciously in the underlying filesystem.

2.7 Hosting Model

There are currently four different methods of configuring a location, three of which serve Shiny applications. These three are described below; see the following section on [Redirecting](#) to learn about a fourth mode.

2.7.1 Host a Directory of Applications

A location which uses `site_dir` will host an entire directory tree -- both Shiny Apps and static assets. This is the location used to serve an asset or application in `/srv/shiny-server/` in the default configuration that ships with Shiny Server.

```
# Define the location '/'
location / {
  site_dir /srv/shiny-server/
}
```

The above configuration would instruct Shiny Server to make the `/srv/shiny-server/` directory available at the base URL. Any Shiny applications stored in this directory (or any subdirectory), along with any static assets (including images, data, JavaScript/CSS files, etc.) will be made available at the corresponding URL. For example, see the following directory tree:

```
+---/srv/shiny-server
|   +---shinyApp1
|       +---server.R
|       +---ui.R
|   +---shinyApp2
|       +---server.R
|       +---ui.R
|   +---assets
|       +---style.css
|       +---script.js
```

If this server were available at `http://server.com`, the location settings above would make the following (along with any other file

in the tree) publicly available to the user.

URL	Definition
http://server.com/shinyApp1	Serve the Shiny App defined in 'shinyApp1'
http://server.com/shinyApp2	Serve the Shiny App defined in 'shinyApp2'
http://server.com/assets/style.css	Serve this static CSS file
http://server.com/assets/script.js	Serve this static JS file

2.7.2 Host a Single Application

A location configured to use `app_dir` will attempt to serve a single application hosted at the given directory. For instance,

```
# Define the location '/specialApp'
location /specialApp {
    app_dir /srv/shiny-server/myApp;
}
```

will configure the location responsible for the path `/specialApp` to use this `app_dir` router which serves a single application stored in the directory `/srv/shiny-server/myApp`. This configuration would presume a `server.R` file (and corresponding `ui.R` file) would be available at `/srv/shiny-server/myApp/server.R`. If so, the application saved there would be available at a URL like `http://server.com/specialApp`.

2.7.3 Host Per-User Application Directories

A location configured to use `user_dirs` will allow users on the system to create and manage their own Shiny applications available in their home directories. This directive will host any application stored in an eligible user's `~/ShinyApps` directory publicly at a URL prefaced by their username.

This privilege can be restricted only to users of particular groups using the `members_of` restriction. For instance, the following configuration:

```
run_as :HOME_USER;;

# Define the root location
location / {
    user_dirs;

    # Only allow members of the 'shinyUsers' group to host personal applications.
    members_of shinyUsers;
}
```

will -- for any user who is a member of the `shinyUsers` group -- publicly host any Shiny application available in the user's `~/ShinyApps` directory. For instance, if `tina` is a user on this system and is also a member of the `shinyUsers` group, any application stored in `/home/tina/ShinyApps/`, such as the `shinyApp1` application, would be available on this server at a URL like `http://server.com/tina/shinyApp1`, assuming `/home/tina` was `tina`'s home directory.

The `user_dirs` setting is intended to ultimately replace the `user_apps` setting. The two models are very similar but differ in how they handle `run_as` settings. `user_apps` ignores the `run_as` setting and always runs applications as the user in whose home directory they exist. `user_dirs`, on the other hand, will respect the `run_as` setting. You can use the special `run_as` user `:HOME_USER:` to make `user_dirs` operate the same way as `user_apps` does; thus, the two following configurations would be identical:

```
location / {
    user_apps;
}
```

```
location / {
    run_as :HOME_USER;;
    user_dirs;
}
```

Note that, while Shiny Server supports the ability to run in `user_dirs` as a single user, this may require some tweaking to make

function correctly on your server. For instance, the following configuration:

```
location / {  
    run_as shiny;  
    user_dirs;  
}
```

would indeed attempt to start the applications housed in your users' home directories as the shiny user, but many distributions, by default, restrict access to users' private /home directories, meaning that the shiny user would not be able to read or write to the directories it needs to in order to run these applications properly.

2.8 Redirecting

The final mode in which a location can operate is to redirect to another URL. Such locations will immediately send a response to the client informing them of the URL to which they should redirect and the status code that should be used when informing the client of the redirection. Typically, a redirect will use the 301 status code for a permanent redirect, or a 302 status code for a temporary redirect. The final option when configuring a location for redirection is whether or not it should use exact matching. If a redirecting location is configured to use exact matching, only requests for that exact URL will be redirected. If not, any requests for that URL path or any subpath of that URL will be redirected. For example,

```
# Define a location at the base URL of this 'server'  
location / {  
    # Redirect traffic from '/shinyApp1/' to 'http://server.com' temporarily.  
    location /shinyApp1 {  
        redirect "http://server.com" 302 true;  
    }  
}
```

will redirect any requests for the exact path /shinyApp1 to http://server.com temporarily.

2.9 Virtual Hosts

The `server_name` setting allows Shiny Server to route incoming requests to a particular server element based on the hostname of the request. This will serve as an additional level of filtering on incoming traffic. Bear in mind that traffic must be destined for this server (i.e. traffic which was bound for the IP address and port on which this server was listening) in order to even be evaluated for a matching hostname.

The example below may clarify the configuration

```
server {  
    # Instruct this server to listen to port 80  
    listen 80;  
  
    # Only accept requests for the hostname 'server1.com'  
    server_name server1.com;  
  
    # Define the location for this server.  
    location / {  
        site_dir /srv/shiny-server1;  
        log_dir /var/log/shiny-server1;  
    }  
}  
  
server {  
    # Instruct this server to listen to port 80  
    listen 80;  
  
    # Only accept requests for the hostname 'server2.com'  
    server_name server2.com;
```

```
# Define the location for this server.
location / {
    site_dir /srv/shiny-server2;
    log_dir /var/log/shiny-server2;
}
}
```

This example presupposes that, on this network, the domains `server1.com` and `server2.com` both resolve to the same IP address. In the configuration above, we first create a server that listens on port 80 and will only accept traffic whose hostname matches `server1.com`. We then configure a server which will also listen on port 80 and require that it only accept traffic whose hostname equals `server2.com`.

This configuration would allow an application stored in `/srv/shiny-server1/shinyApp` to be accessible at `http://server1.com/shinyApp` and an application stored in `/srv/shiny-server2/shinyApp` to be accessible at `http://server2.com/shinyApp`.

2.10 Custom Templates

Shiny Server can utilize custom templates when generating static pages such as directory indexes and error pages. This feature is controlled using the `template_dir` directive and can be applied globally, or to a particular server or location.

To utilize this feature, you'll need a directory in which you'll store your templates. In this example, we'll assume you're using `/etc/shiny-server/templates/`. Inside this directory, you'll place the `handlebars`-enabled HTML templates for your pages.

Shiny Server will attempt to find a specific page to accommodate the current need -- for instance, if it encountered a 404 error, it will first look to see if you provided a specific template for handling 404 errors in the specified directory (`error-404.html`). If it doesn't find this file, it will go one step backwards to see if you provide a generic error template (`error.html`). If it can't find a suitable template in your template directory, it will fall back to use the provided templates that come with Shiny Server (and are stored in `/opt/shiny-server/templates`).

As you may have noticed already, Shiny Server uses a hyphen-delimited hierarchy which progresses from broadest to narrowest. So all error pages start with `error`. But the file name for a 404 error page would add an additional element (`404`) following a hyphen: `error-404`. All files are expected to use the `.html` extension.

The tables below describe the `page templates that can be defined in Shiny Server`:

Template Name	Description
error-403	Client is forbidden from accessing this page. Or, in Shiny Server Pro, that the user is signed in but does not have permissions to view this application.
error-404	Page not found
error-500	General server error
error-503-users	The application has exceeded its maximum # of users
error-503-license	(Pro Only) You have exceeded the number of concurrent users allotted for your license.
directoryIndex	Shown when <code>directory_index</code> is enabled for a location and the user visits the base URL.
login	(Pro Only) Used to render the login page.

Again, any of the template names that have hyphens in them can be generalized by trimming off the hyphen and subsequent word(s). So a file named `error-503.html` could be used as the template for *both* `error-503-users` pages, as well as `error-503-license` pages. More generally, a single template can be used for all errors by naming it `error.html`. If a more specific template doesn't exist, this page would be used when generating any type of error page.

We recommend starting with the existing templates then modifying them to create your own branding. To do this, copy the relevant HTML files out of `/opt/shiny-server/templates/` into your template directory (for example, `/etc/shiny-server/templates/`). To create a custom error page, execute:

```
$ mkdir /etc/shiny-server/templates/
$ cp /opt/shiny-server/templates/error.html /etc/shiny-server/templates/
```

Now open the template in your preferred text editor and modify it. Initially, leave the overall structure intact, but add some

language at the top of the page (immediately under the line containing `<h1>{{title}}</h1>`) or add some new CSS rules then save the file.

Now instruct Shiny Server to use this new template in your configuration by adding the line:

```
template_dir /etc/shiny-server/templates;
```

at the top/bottom of the file so it applies globally. Of course, you can later specify another `template_dir` for particular locations.

When you restart the server and access an error page. (Try accessing a hidden file like `http://server.com:3838/.hidden`, where `server.com` is the domain name or IP address of your server, which should use the `error-403` or `error` template.) You should see your new text or CSS styling on the page that is loaded. Likewise, you can copy the `directoryIndex.html` template or, for Shiny Server Pro, the `login.html` template to your custom directory and begin customizing.

For performance reasons, templates are cached in memory once they're loaded. Thus, you will find that changes to your template may not be immediately visible when you refresh the page. To empty the cache to force Shiny Server to load in the new changes to your template, you can either restart or reload the server (see the section on [Stopping and Starting](#)).

Finally, note that these templates currently only apply for HTML pages that are generated by Shiny Server. An HTML page generated by the Shiny package will *not* have custom templates applied at this time. Currently, this is most visible with 404 errors. If you access a URL which Shiny Server cannot direct to a Shiny application (for instance, `http://server.com:3838/flargdarg`, where `server.com` is the domain name or IP address of your server), you will see a Shiny-Server-generated 404 error page which has custom templates applied (assuming this server is configured to have a `custom-error-404.html` or `error.html` template). However, if you were to access a file that does not exist in a valid application, you would see an un-styled 404 page that did not leverage your custom template. For instance, if you have an application deployed at `http://server.com:3838/myApp` and you attempt to access `http://server.com:3838/myApp/flargDarg`, the request would be sent to Shiny which would then look for this file and, if it isn't found, would return a 404 page. Shiny Server would pass this response back directly to the client, as it wasn't generated by Shiny Server. We will be reviewing this architecture in future versions -- please let us know if you find that the current organization inhibits your workflow.

2.11 Server Log

All information related to Shiny Server, rather than a particular Shiny application, is logged in the global system log stored in `/var/log/shiny-server.log`. This log should be checked often to ensure Shiny Server is performing as expected. Any errors and warnings which Shiny Server needs to communicate will be written here.

If `logrotate` is available when Shiny Server is installed, a `logrotate` configuration will be installed. The default configuration is to rotate the logfile when it exceeds 1MB. The old log file will be compressed and stored alongside the original log file with a `.1.gz` extension (then `.2.gz`, etc.). Up to twelve archived log files will be maintained; upon the 13th log rotation the oldest log file will be deleted.

For logs related to individual Shiny Applications, see the section on [Logging and Analytics](#).

2.11.1 Access Logs

An access log can be configured globally using the `access_log` parameter. This log is not enabled by default. This setting controls the location of the access log as well as the format used. The access log can be useful to audit the security and activity taking place on your Shiny Server installation. These logs can be audited manually or automatically to inspect how often various resources are being accessed, or by whom they are being accessed (using the originating IP address). Currently, one access log is created for the entire Shiny Server process; all Shiny applications share this access log.

The access logs will be written using the Connect logging libraries, so additional details can be found in [their documentation](#) in the "Formats" section. In brief, there are four pre-defined logging formats for access logs specified in the documentation referenced above:

- default `'remote-addr - - [:date] ":method :url HTTP/:http-version" :status :res[content-length] ":referrer" ":user-agent"`
- short `'remote-addr - :method :url HTTP/:http-version :status :res[content-length] - :response-time ms'`
- tiny `'method :url :status :res[content-length] - :response-time ms'`
- dev concise output colored by response status for development use

For example, the following would configure an access log that uses the tiny format:

```
access_log /var/log/shiny-server/access.log tiny;

server {
```

2.12 Environment Variable Settings

We strive to expose most settings through the configuration file described previously. However, there are some settings that we don't include in the configuration file which some users may still wish to change. Where possible, we'll allow users to configure these settings using environment variables. Some of these environment variables will later be promoted to configuration directives, but will be documented here until that time.

Typically, it is best to define these environment variables in the startup script used to run Shiny Server. This varies depending on what Linux distribution and version you are running.

2.12.0.1 systemd (RedHat 7, Ubuntu 15.04+, SLES 12+)

File to change:

```
/etc/systemd/system/shiny-server.service
```

How to define the environment variable:

```
[Service]
Environment="SHINY_LOG_LEVEL=TRACE"
```

Commands to run for the changes to take effect:

```
$ sudo systemctl stop shiny-server
$ sudo systemctl daemon-reload
$ sudo systemctl start shiny-server
```

2.12.0.2 Upstart (Ubuntu 12.04 through 14.10 and RedHat 6)

File to change:

```
/etc/init/shiny-server.conf
```

How to define the environment variable:

```
env SHINY_LOG_LEVEL=TRACE
```

Commands to run for the changes to take effect:

```
$ sudo restart shiny-server
```

2.12.0.3 init.d (RedHat 5, SLES 11)

File to change:

```
/etc/init.d/shiny-server
```

How to define the environment variable:

```
export SHINY_LOG_LEVEL=TRACE
```

Commands to run for the changes to take effect:

```
$ sudo /sbin/service shiny-server restart
```

2.12.1 SHINY_LOG_LEVEL

Defines the verbosity of logging which will be used by Shiny Server. Valid options -- in level of increasing verbosity -- are TRACE, DEBUG, INFO, WARN, and ERROR. More verbose levels of logging (such as TRACE) may be helpful when debugging an issue or trying to understand exactly what Shiny Server is doing, but will likely be far too much information for a system with even moderate load.

The default if this environment variable is not found is to use INFO.

2.12.2 R

Defines the path to the R executable which should be used when running Shiny. Systems which have multiple versions of R installed, or don't wish to place R on the path before starting Shiny Server can override this setting to point to a particular version of R.

Note that, for any location with a configured `[r_path]` in the configuration file, this environment variable will not take effect. If no environment variable is found, Shiny Server will expect an executable named R to be on the path.

2.12.3 SHINY_DATA_DIR - Pro Only!

The historical databases of Shiny applications used in the Admin dashboard can consume multiple Gigabytes of disk space on a system hosting many applications. This environment variable allows you to define the directory in which Shiny Server should persist data to disk. To estimate the amount of disk space required, we recommend reserving 10MB in this directory for every Shiny application you plan to host on a server. Additionally, ensure that this directory is writable by the shiny user, or whichever user you're running Shiny Server as (as is discussed in detail in the [run_as](#) section).

Note that Shiny Server expects a directory named `monitor/rrd/` to already exist inside `SHINY_DATA_DIR` when it starts.

By default, Shiny Server will create and store data in `/var/lib/shiny-server/`.

3 Deploying Applications

3.1 Schedulers & Application Restarts

A scheduler is responsible for fulfilling incoming requests to a particular application. Each version of an application (see below) will have its own associated scheduler. Each scheduler can have different properties to control things like how many concurrent connections it should accept.

A scheduler can be specified at many locations in the configuration file and will be inherited to inner blocks. For instance, a scheduler definition found in a server block will be applied to every location in that server, unless overridden.

3.1.1 Restarting an Application

Some changes you make to the code, assets, or environment of an application will require the application's R processes to be restarted for the changes to take effect. These include upgrades to packages that are used by the application, changes to `.Renv` or `.Rprofile` or other R source files, or modifications to data files that are read only at startup time.

(Fortunately, Shiny applications generally do not need to be restarted when changes to `ui.R` or `server.R` are made, as Shiny will check for changes to these files on page load.)

An application can be restarted by altering the "modified time" on a file named `restart.txt` in the application's directory. This can most easily be done using the `touch` utility as in `touch restart.txt` which will update the modified timestamp on this file to the current time. Upon the next new connection to the application, Shiny Server will spawn a new R process to run the "new" (restarted) Shiny Application for this and future users. When this occurs, the old processes will remain unaltered, and open connections will remain active and valid until the last connection closes itself.

This can cause unintuitive consequences regarding the number of processes that may be running for a given application. Even if using a single-process Simple Scheduler for an application, it is possible that there would be multiple R processes associated with this application, each corresponding to a separate restart timestamp. This behavior is likely to change in future versions of Shiny Server.

3.1.2 Simple Scheduler

The Simple Scheduler is the only scheduler available in the Open Source edition of Shiny Server. It associates a single R process with a single Shiny application. This scheduler accepts a single parameter which specifies the maximum number of concurrent

sessions. Once this number is reached, users attempting to create a new session on this application will receive a 503 error page.

```
# Define a new location at the base URL
location / {
  # Define the scheduler to use for this location
  simple_scheduler 15;

  ...
}
```

In this example, the location defined at the base URL in this server will be configured to use the Simple Scheduler and to limit each application to a maximum of 15 simultaneous connections -- meaning that if there are 15 active connections on the application when a user attempts to visit this application, that user will receive a 503 error page. When one of the 15 active connections is disconnected, one new "seat" will be available for another user to connect to the application.

3.1.3 Utilization Scheduler - Pro Only!

The Utilization Scheduler is a more sophisticated scheduler only available in Shiny Server Professional that allows a single Shiny application to be powered by multiple R Shiny processes. Incoming traffic will be routed to the R process associated with this Shiny application which has the fewest current connections.

The Utilization Scheduler is configured by three parameters.

- `maxRequestsPerProc` -- The number of concurrent connections which a single R process can support before it will begin to return 503 errors for new sessions.
- `loadFactor` -- The "trigger percentage" (on a scale from 0 to 1) at which point a new R process should be spawned. Once session capacity reaches this percentage, a new process will be spawned (unless the `maxProc` limit has been reached).
- `maxProc` -- The maximum number of processes which should exist concurrently associated with this Shiny application. Note that this does not include restarted versions of the same application in its count.

As an example:

```
location /shinyApp1 {
  # Define the scheduler to use for this application
  utilization_scheduler 5 0.5 3;

  ...
}
```

The above configuration would create a Utilization Scheduler for this application which would support a maximum of 3 processes, a maximum of 5 connections per process, and a load factor of 0.5. Initially, only one process would be created to facilitate the first and second concurrent connection to this application. The third connection, however, would surpass the load factor ($3/5 > 0.5$) which would spawn a new R process. The next request would then be routed to this new process. This would continue until the load factor was again surpassed, at which point the third process would be spawned. With a maximum of 3 processes, no more would be created after the third. So the 16th concurrent connection would be turned away with a 503 error.

Be aware that, upon the restart of an application using a Utilization Scheduler, the resources allotted to this application will effectively double for a time. For instance, if an application is configured to use 4 `maxProc` and has sufficient load to create and maintain that many R processes, a restart of the application will motivate Shiny Server to create an R process for a "new" Shiny application which also is eligible to run up to 4 Shiny processes in parallel. Thus, with one restart and sufficient traffic to this application, it is possible that it could be running 8 Shiny processes. This behavior is likely to change in future versions of Shiny Server.

3.2 R Markdown

In addition to serving traditional Shiny applications out of a directory (a `server.R` file and an associated UI), Shiny Server now supports interactive R Markdown documents. If a hosted directory does not include a `server.R` file, Shiny Server will look to see if it contains any `.Rmd` files. If so, Shiny Server will host that directory in "R Markdown" mode using `rmarkdown::run`.

Particular `Rmd` files can be accessed by referencing their full path including the filename, e.g.

<http://myserver.org/mydocs/hello.Rmd>. If a request is made to a directory rather than a particular `Rmd` file, Shiny Server will attempt to serve the `file index.Rmd`. If that file does not exist, the user will get an error alerting them that the file is not available.

3.3 Application Timeouts

Each Shiny Application has two timeouts associated with it:

1. `app_init_timeout` -- Describes the amount of time (in seconds) to wait for an application to start. After this many seconds if the R process still has not become responsive, it will be deemed an unsuccessful startup and the connection will be closed.
2. `app_idle_timeout` -- Defines the amount of time (in seconds) an R process with no active connections should remain open. After the last connection disconnects from an R process, this timer will start and, after the specified number of seconds, if no new connections have been created, the R process will be killed.

Typically, these two parameters will be correlated. Shiny Applications which involve little processing to start (therefore have a small `app_init_timeout`) can often be closed with minimal concern (thus would have a small `app_idle_timeout`). Conversely, applications which require a substantial amount of data to be loaded on startup may merit a longer `app_init_timeout` to give the data time to load, and a longer `app_idle_timeout` as the task of spawning a new process is more expensive and should be minimized.

3.4 Session Timeouts - Pro Only!

The `app_session_timeout` setting can be used to automatically disconnect idle Shiny connections. The setting accepts a single parameter which dictates the number of seconds after which an idle session will be disconnected. Here "idleness" is measured by a connection's interaction with the server. Incoming or outgoing data will reset the countdown time for this user's session. If the number of seconds specified elapses without any data being sent to or from this client, the user's session will be disconnected.

This parameter can be configured as a setting globally, or of a particular server or location. It can be used to ensure that stale R processes are eventually closed by terminating sessions after they become idle for some period of time.

Session timeouts are configurable on a per-server or per-application basis. Controlling the session timeout can be important in helping the administrator manage the resource allocation for a given application and in freeing up resources from idle connections which will then be available to newly created sessions.

3.5 Logging and Analytics

3.5.1 Application Error Logs

Error logs are created for each R Shiny process separately. For locations configured to use `user_apps` or `user_dirs`, these logs are created in each user's `~/ShinyApps/log/` directory. For locations configured with `app_dir` or `site_dir`, the directory in which these logs are created is managed by the `log_dir` setting, which can be specified globally, for a particular server, or for a particular location. By default, these logs will be located at `/var/log/shiny-server/`. To change these locations, the `log_dir` can be specified in the following way:

```
location / {  
  log_dir /var/log/shiny-server/;  
}
```

The log files will be created in the following format:

```
<application directory name>-YYMMDD-HH:mm:ss-<port number or socket ID>.log
```

Log files will be created for each R process when it is started. However, if a process closes successfully, the error log associated with that process will be automatically deleted. The only error log files that will remain on disk are those associated with R processes that did not exit as expected.

If you're looking for log messages related to Shiny Server itself -- rather than individual Shiny applications -- see the section on the [Server Log](#).

3.5.2 Google Analytics

Shiny Server is capable of automatically inserting the necessary JavaScript code to enable Google Analytics tracking globally or for a particular server or location. This is managed by the `google_analytics_id` setting which can be used as follows:

```
location / {
```



```
google_analytics_id UA-12345-1;  
}
```

3.6 Program Supervisors - Pro Only!

You may also wish to run R sessions under a program supervisor that modifies their environment or available resources. You can specify a supervisor (and the arguments which control it's behavior) using the `exec_supervisor` setting. For example:

```
...  
  
location /low {  
  
    ...  
  
    exec_supervisor "nice -n 10";  
  
}
```

This example uses the `nice` command to run the Shiny processes in the affected location with a lower scheduling priority. See <http://linux.die.net/man/1/nice> for more details on `nice`. Note that for `nice`, in particular, it's possible to accomplish the same thing using the `pam_limits` module (and even specify a custom priority level per user or group).

It is important to note that the `exec_supervisor` will be executed as the same user the Shiny app will be running as. So if Shiny Server is running as root but runs applications as `shiny`, the `exec_supervisor` will be executed as `shiny`. Because most systems only allow the root user to assign a negative (higher priority) `nice` privilege, this means that `exec_supervisor` would not be able to assign a negative `nice` value without reconfiguring the server to allow certain users to assign negative `nice` values to a process.

3.7 Memory Limits - Pro Only!

Memory limits will be added in a coming version of Shiny Server Professional.

3.8 Reactivity Log

The reactivity log is a great browser-based tool for analyzing and debugging reactive dependencies in a Shiny Application. We strongly recommend that you read the [Overview of Reactivity](#) in the Shiny tutorial to get the most out of the reactivity log. The symbols and representations introduced in that article are reused in the log. This tool should not be used in a production environment, as it exposes details that should not be publicly available. Many application authors find it to be a useful tool to use locally, however.

Currently, reactivity logs are maintained for an entire R process. Therefore, it should be used in a fresh R session in which only one Shiny application is active. To enable the reactivity log, run the command `options(shiny.reactlog=TRUE)` before running your application. Once you've started your application and are viewing it in a browser, you can use `Ctrl+F3` (or `Command+F3` for Mac users) to view the reactivity log. This visualization creates a snapshot of all reactive activity that had occurred up to the moment it was created. If you want to view the reactive activity that has occurred since opening the reactivity log, you'll need to refresh the browser.

The `showReactLog()` function exposes this functionality from within R and can be used outside of strictly Shiny contexts to generate a static HTML file visualizing reactivity. You can see the full documentation using `?showReactLog` in R.

As the reactivity log exposes the architecture and some of the code behind your application, it would likely be unwise to enable this feature in a production environment. Unless explicitly enabled by the user, Shiny Server will not enable this functionality. We recommend using the reactivity log in a local R process not managed by Shiny Server.

3.9 Specifying Protocols

Shiny Server can use a wide variety of techniques to keep the data in the web browser synchronized. The preferred technique, and the one most widely used, is the use of WebSockets. However, if WebSockets are not supported -- either by some intermediate network between Shiny Server and your client or by your client's web browser -- then a fallback protocol will be

used. In total, Shiny Server provides nine different methods of connecting to the server in real-time. In order of preference, they are:

1. WebSocket
2. XDR Streaming
3. XHR Streaming
4. iframe Eventsource
5. iframe HTML File
6. XDR Polling
7. XHR Polling
8. iframe XHR Polling
9. JSONP Polling

Each of these will be tried by the client in order until a successful connection to Shiny Server. If you would like to omit one or more of these protocols from the list for whatever reason, you can currently do so only from the client (we hope to make a server-side configuration available in a future release of Shiny Server).

To change the available protocols from the client, you can open any Shiny Application and press the keyboard shortcut: Ctrl+Alt+Shift+A (or, from a Mac: control+option+shift+A). This will bring up a window that will allow you to select or deselect any of the above protocols. After confirming your changes these settings will be saved in your browser for future visits to this server. These settings will take effect upon loading any Shiny application hosted on this domain and will last until you explicitly change them again; they will only have an effect on the browser in which this action was performed.

3.9.1 Disabling WebSockets on the Server

Shiny Server uses SockJS to facilitate communication between the browser and the server. On most networks, SockJS will be able to resolve a functioning protocol that will allow it to reliably communicate between the browser and Shiny Server. Some environments, however, will perform some sort of filtering or proxying that hinders one or more of the available protocols. Older networking equipment will likely not function well with WebSockets, for instance.

Most users will not need to worry about what protocol their network will support, as SockJS will gracefully determine the best protocol when trying to establish the connection. However, some environments have found consistent problems with having particular protocols enabled and find that a particular protocol or set of protocols are most reliable in their environment. In this case, you can use the `disable_protocols` setting to disable any of the available protocols.

For instance, if you find that WebSockets and XDR polling are problematic on your networking equipment, you could add the following configuration to disable them:

```
disable_protocols websocket xdr-polling;
```

The protocol names to be used when disabling are as follows:

```
websocket
xdr-streaming
xhr-streaming
iframe-eventsource
iframe-htmfile
xdr-polling
xhr-polling
iframe-xhr-polling
jsonp-polling
```

4 Authentication & Security - Pro Only!

4.1 Authentication Overview - Pro Only!

Shiny Server Professional offers the ability to authenticate individual users. By specifying authentication requirements on particular servers or locations, the administrator can control the set of applications particular users are allowed to access.

Authentication can be implemented by integrating into an existing LDAP or Active Directory database, relying on Google accounts, or by using a "flat-file" authentication system which is simply contained in a local file. To ensure that your users' passwords are being protected, it is strongly encouraged that any Shiny application requiring authentication use [SSL](#) to encrypt the usernames and passwords being transmitted.

The current user's username and groups (where applicable) will be available in Shiny (version 0.8 and later) in the session parameter of the shinyServer function. For instance:

```
shinyServer(function(input, output, session) {  
  output$username <- reactive({  
    session$user  
  })  
  
  output$groups <- reactive({  
    session$groups  
  })  
})
```

It is important to note that when a user logs in to Shiny Server Pro, that authentication will be applied server-wide. This means that a visitor who authenticates as a particular user on one application will be identified as that user to the other applications on this server to which they connect. There is currently no notion in Shiny Server Pro of logging in exclusively to only one particular application.

4.2 Auth Duration - Pro Only!

Regardless of which authentication mechanism is used, the duration of authentication can be configured using the `auth_duration` setting. This setting controls the amount of time (in minutes) for which a user should remain logged in after the finish using their last application. This setting should be set to 10 minutes or greater.

4.3 Flat-File Authentication - Pro Only!

Shiny Server Pro offers flat-file authentication as a simple and easy-to-configure authentication mechanism. This method is self-contained and not integrated into either the system's user/password database, nor any Enterprise authentication mechanism. Thus, usernames and passwords must be created explicitly for each user that should exist in Shiny Server.

The storage of usernames and passwords is handled by a single file which can be specified using the `auth_passwd_file` setting as follows:

```
run_as shiny;  
  
auth_passwd_file /etc/shiny-server/passwd;  
  
server {  
  location / {  
    ...  
  }  
}
```

This will instruct Shiny Server to lookup all usernames and passwords in the file stored at `/etc/shiny-server/passwd`. This file should have zero or more lines in the format `username:{scrypt-hashed-password}`. The scrypt encryption algorithm is used to protect users' passwords from theft; the hashed passwords are expected in this file in base64 format. We provide the [sspasswd](#) utility with Shiny Server Professional that vastly simplifies the process of managing these "sspasswd" files.

You'll want to think carefully before adjusting the permissions on this file. For instance, if you were to enable the shiny user to write to your password database, any Shiny application running as shiny (the default) would now be able to modify your password database. Because the passwords are securely hashed, granting shiny read access to this file is not problematic and, in fact, is enabled by default.

4.3.1 sspasswd - Pro Only!

The `sspasswd` utility comes with Shiny Server Professional and can be used to manage the username/password file. By default, it is not made available on the PATH but you can find it in `opt/shiny-server/bin/`. The general pattern for the utility is to provide the file to use for storage followed by a username, as in:

```
$ sudo /opt/shiny-server/bin/sspasswd /etc/shiny-server/passwd tina
```

To create a (or overwrite an existing) password file, use the `-c` switch. The default behavior will be to add the username specified

(tina, in the example above) to the file after prompting the user for a password (or reading it from stdin). To instead delete a user from the file, use the -D switch. User/password combinations can be verified using the -v switch. Finally, the -C switch will set the maximum amount of time (in seconds) for encryption. The larger this value is, the more secure the password hash will be.

Currently, the ':', '\$' and newline ('\r' or '\n') characters are prohibited in usernames.

4.4 Google Authentication - Pro Only!

Shiny Server Pro supports the ability to rely on Google for the management of your users using the `auth_google` setting. In this environment, users attempting to login to your Shiny Server would be presented with a button asking them to login via Google. If they grant access to your application, they would be returned to your Shiny Server presenting the email address associated with their Google account. At that point, this email address can be used to grant or restrict access to individual applications and locations. This section will show you how to setup this type of Google Authentication for Shiny Server.

4.4.1 Create a Google Application - Pro Only!

Before you can use Google Authentication on your server, you must register an "application" with Google. This is the application that your users will be asked to "connect" to when logging in to your Shiny Server. To do this, you can visit Google's Developer's Console which, at the time of writing, is available at <https://console.developers.google.com/>.

You'll initially create a project then give it a name and ID. Under "APIs & auth" you'll find the "Credentials" section. There you have the ability to "Create new Client ID"; do that now. The type of application is a "Web application". For the "Authorized JavaScript Origins", you'll need to explicitly list any name a user might use to access your server including the domain name, port, and protocol. For instance, if you offer both HTTPS and HTTP access to your Shiny Server, you'll need to include both; if you have multiple domain names pointing to your server, you should include all of those domains; if you have servers running on multiple ports, you'll need to include them all explicitly here.

In the "Authorized redirect URI" field, you'll again need to provide all possible ways a user may be accessing your server, but append the string `__login__/callback` (note the double underscores on either side of "login") to each URI.

For instance, if you have a server available at both `example.org` and `example.net`, supporting both HTTP and HTTPS, offering web services with or without a `www.` subdomain prefix, and using the standard ports (80 and 443, respectively), you would use the following origins:

- `http://example.org`
- `http://www.example.org`
- `https://example.org`
- `https://www.example.org`
- `http://example.net`
- `http://www.example.net`
- `https://example.net`
- `https://www.example.net`

And the Authorized redirect URIs would include:

- `http://example.org/__login__/callback`
- `http://www.example.org/__login__/callback`
- `https://example.org/__login__/callback`
- `https://www.example.org/__login__/callback`
- `http://example.net/__login__/callback`
- `http://www.example.net/__login__/callback`
- `https://example.net/__login__/callback`
- `https://www.example.net/__login__/callback`

You can then "Create" your Client ID.

You should see your new "Client ID for web application" listed once you have created your new Client ID. You can inspect to confirm that the Redirect URIs and Javascript Origins look correct (if not, you can click "Edit settings"). If they look correct, you should note the "Client ID" and "Client secret"; you'll enter these settings into your Shiny Server configuration file later.

At this point, you've created a Google application you'll be able to use to allow your users to login via Google on your Shiny Server. You can always add more details to your Google application here including an icon and more textual details.

4.4.2 Configure Shiny Server - Pro Only!

Now that you have a Google application created, you can customize your Shiny Server configuration to instruct it to a.) Use Google for its authentication, and b.) provide your Google application's settings.

You can edit your configuration file stored at `/etc/shiny-server/shiny-server.conf` to replace any existing authentication system (if you have one) with the `auth_google` directive. This setting first takes the Client ID, then the Client Secret you created in Google earlier like so:

```
auth_google 12345.google.com ABCDEF;
```

At this point, you can restart your Shiny Server (see [Stopping and Starting](#) if you're unsure how). When you try to access a protected resource such as the Admin page or a location secured with a `required_user` setting, you should be prompted with a login screen asking you to login with Google. Click that link, then decide whether or not you'd like to grant access to your new application. If you do grant access and your email is included in the `required_user` field associated with the resource you're trying to access, you'll be let in. Otherwise, you'll still be logged in as that user, but won't be granted access to that resource.

Do be aware that it's a good idea to protect your OAuth application secret from unprivileged users. If you're using `auth_google`, it is strongly encouraged that you read the following section on [Securing the Client Secret](#).

It may be necessary to educate your users on the behavior of Single Sign On (SSO) systems like Google. Specifically, some users may not understand that logging out from a Shiny app will *not* log them out of Google. Currently, logging out of Google will also not log you out of Shiny Server. Rather, there is a one-time connection between the two systems at the moment the user logs in. From there on, their Google session and their Shiny Server session will be completely independent.

4.4.3 Securing the Client Secret - Pro Only!

In addition to providing the client secret as text in the config file, you could alternatively specify the absolute path to a file containing the client secret as in the following example:

```
auth_google 12345.google.com /etc/shiny-server/ga-secret.txt;
```

In this case, Shiny Server would read the client secret in from the file stored at `/etc/shiny-server/ga-secret.txt`. Additionally, if the client secret is a path to a file rather than the secret itself, Shiny Server will retain root privileges, allowing you to tighten the security around the file containing your client secret to allow only the root user to read it. For example, we could secure the above file using:

```
# Change the file to be owned by the root user. {.pro}
chown root:root /etc/shiny-server/ga-secret.txt

# Set the permissions so that only root can read the file. {.pro}
chmod 600 /etc/shiny-server/ga-secret.txt
```

In this way, you can secure your client secret from Shiny applications or other users who are allowed to login to the server. If you expect that many users or untrusted users may be logging into or deploying apps on the server on which Shiny Server is installed, we strongly recommend that you use this configuration to ensure the privacy of your secret key.

4.4.4 Email Address Restrictions - Pro Only!

By default, the Google auth strategy will allow any user with a Google account to login to your system. Of course, the `required_user` settings would continue to secure your resources and only provide access to resources the user is authorized to see.

Though you can always use `required_user` to restrict access to a resource, some organizations wish to only allow particular users to login at all. For instance, if you were not concerned about which individual users could access which applications, but wanted to ensure that only users within your organization can access *any* application, you may want a more broad restriction on the users you authorize. As discussed in the section on [User Authentication](#), you can use `required_user *`; to indicate that any logged in user should be able to access a resource.

For such needs, we provide the ability to globally filter which email addresses will be authorized into your Shiny Server at all -- regardless of what the individual application's `required_user` settings are. To use this, you can add filters onto the `auth_google` setting in your configuration file. In the simplest form, you could add one filter to require that the email address the user is using with Google be of a particular form, such as `"*@myorganization.org"`. This would allow any user who logged into Google with the email address that ended in `@myorganization.org` to log on to the server. Once logged on, the `required_user` filters would be applied as usual to control access to individual applications.

Additionally, you can provide as many filters as you'd like on the `auth_google` directive and indicate whether email address matching that filter should be admitted (+) or rejected (-). They will be applied one-at-a-time until the given email address matches one of your filters. The filters can be prefixed with a + or - sign to indicate whether the filter is "positive" (in which case any user matching the filter should be admitted) or "negative" (in which case any user matching the filter should be rejected). If

no sign is provided, it's assumed that the filter is "positive". Filters will be evaluated *in the order in which they are provided*, and the first filter an email matches will be used to determine whether or not that email address should be allowed to login to Shiny Server. If an email address matches none of the provided filters, it will be rejected. A few examples may be helpful:

```
auth_google 12345.google.com ABCDEF "*";
```

The filter "*" will be interpreted as a positive filter, meaning any email address matching that filter (every possible email address) will be admitted. This is the default behavior if you don't provide any filters in the configuration file.

```
auth_google 12345.google.com ABCDEF "@myorganization.org";
```

The filter "@myorganization.org" will admit any user whose email address ends in @myorganization.org and reject any email address that does not.

```
auth_google 12345.google.com ABCDEF "-jim@example.org" "@example.org";
```

In these filters, if the user kelly@example.org were attempting to the system, the first filter would be checked and it would be found that kelly@example.org does not match the filter "jim@example.org", so it would continue on to the next filter. There it would be found that kelly@example.org *does* match the filter *@example.org, and because this filter is a positive filter (you could have prefixed it with a +, if you wished), this user would be granted access to the system. However, a user john@gmail.com would not match the first filter or the second filter, so he would be denied access to the system, as he did not match any positive filter. Finally, the user jim@example.org would match the first filter, and because that filter is a negative filter, this user would be denied access to the system.

We support a lightweight pattern matching syntax. Specifically, you can use the * character to match any string of length 0 or more, the ? character will match any *one* character, and square brackets can be used to match against a set of characters such as [a-e] to match any character between a and e. All patterns are *not* case-sensitive.

4.5 PAM Authentication - Pro Only!

Shiny Server Professional can authenticate users via the Linux standard PAM (Pluggable Authentication Module) API. PAM is typically configured by default to authenticate against the system user database (/etc/passwd) however it can also be configured to authenticate against a wide variety of other systems including Active Directory and LDAP.

This section describes the PAM configuration used for authentication. Note that PAM can be used for both authentication as well as to tailor the environment for user sessions (PAM sessions). This section describes only authentication, see the [PAM Sessions](#) section for details on how Shiny Server can be configured to use PAM sessions.

4.5.1 PAM Basics - Pro Only!

PAM profiles are located in the /etc/pam.d directory. Each application can have its own profile, and there is also a default profile used for applications without one (the default profile is handled differently depending on which version of Linux you are running).

To learn more about PAM and the many options and modules available for it see the following:

- http://en.wikipedia.org/wiki/Pluggable_authentication_module
- http://www.centos.org/docs/5/html/Deployment_Guide-en-US/ch-pam.html
- <http://tldp.org/HOWTO/User-Authentication-HOWTO/x115.html>
- <http://linux.die.net/man/8/pam>

4.5.2 Default PAM Configuration - Pro Only!

Debian/Ubuntu

On Debian and Ubuntu systems Shiny Server does not provide a Shiny-Server-specific PAM configuration file. As a result, Shiny Server uses the /etc/pam.d/other profile, which by default inherits from a set of common configuration files:

/etc/pam.d/other

```
@include common-auth
@include common-account
@include common-password
```

```
@include common-session
```

If the `/etc/pam.d/other` profile reflects the authentication system and policies that you'd like Shiny Server to use then no further configuration is required. If you want to create a custom PAM profile for Shiny Server you would create a file named `/etc/pam.d/shiny-server` and specify whatever settings are appropriate.

RedHat/CentOS/SLES

On RedHat and CentOS systems, applications without their own PAM profiles are denied access by default. Therefore to ensure that Shiny Server is running and available after installation a default PAM profile is installed at `/etc/pam.d/shiny-server`. This profile is configured to require a user-id greater than 500 and to authenticate users against local system accounts:

/etc/pam.d/shiny-server

```
auth      requisite    pam_succeed_if.so uid >= 500 quiet
auth      required     pam_unix.so nodelay
account   required     pam_unix.so
```

This default PAM profile may not reflect the authentication behavior that you want for Shiny Server. In that case, some customization may be required. If you've already set up another PAM profile (e.g. `/etc/pam.d/login`) with the desired behavior then it may be enough to simply copy that profile over the Shiny Server one. For example:

```
$ sudo cp /etc/pam.d/login /etc/pam.d/shiny-server
```

4.5.3 PAM and Kerberos

Shiny Server Professional supports integration with Kerberos for seamless authentication to other applications via Kerberos tickets. To enable this feature you need to have the following defined in your `/etc/shiny-server/shiny-server.conf` file:

```
run_as :AUTH_USER;;
auth_pam true;
pam_sessions_profile shiny-session;
```

This assumes you are using a file named `shiny-session` for the `pam.d` session profile, but the name could be anything as long as it matches your actual filename.

You also need to customize your Shiny Server `pam.d` files: `shiny-server` and `shiny-session`. We provide a very simple example of these files here that uses `pam_krb5.so` as a guide.

/etc/pam.d/shiny-server

```
auth      sufficient    pam_krb5.so
account   required     pam_krb5.so
session   requisite     pam_krb5.so
```

/etc/pam.d/shiny-session

```
auth      required      pam_krb5.so
account   [default=bad success=ok user_unknown=ignore] pam_krb5.so
password  sufficient     pam_krb5.so use_authtok
session   requisite     pam_krb5.so
```

4.6 LDAP and Active Directory - Pro Only!

Lightweight Directory Access Protocol (LDAP) is a popular protocol for storing and validating user information in an enterprise. LDAP can store information about users and their group memberships which Shiny Server Pro is able to query using user's username and password. Fundamentally, the LDAP protocol simply describes a framework for storing hierarchical data and can be configured in a multitude of ways. Active Directory is one popular directory service which implements LDAP and encourages a certain model for storing data, while other vendors of LDAP systems often have their own distinct default configurations. A

holistic overview of LDAP is outside of the scope of this document, so if you lack a solid background in LDAP you might benefit from consulting with an LDAP administrator in your organization to configure these settings.

An LDAP configuration in Shiny Server Pro might look like this:

```
auth_ldap ldap://ldap.example.org/dc=example,dc=org {  
  group_filter "memberUid={username}";  
  group_search_base ou=Groups;  
}
```

It is fully possible to use the `auth_ldap` configuration to integrate with an Active Directory system. However, Active Directory encourages a certain pattern which we capture in the `auth_active_dir` directive which extends `auth_ldap`. The hope is that this configuration can be used to make setup simpler for groups using Active Directory. If you're using an Active Directory server, you may just be able to specify the `auth_active_dir` setting and not worry about providing any of the others.

An Active Directory configuration setting may look something like the following:

```
auth_active_dir ldaps://dc01.example.org/dc=example,dc=org example.org{  
  trusted_ca /etc/ssl/certs/example-org.cert;  
}
```

We currently support only a limited subset of Unicode characters as usernames for LDAP and Active Directory. Usernames must include only:

- Alphanumeric characters
- `_` - Underscore
- `.` - Period
- `@` - "at" symbol

We do not permit empty usernames or passwords.

The parent directive for all LDAP-related settings is `auth_ldap` or `auth_active_dir`. Both accept an LDAP URL as their first argument. `auth_active_dir` accepts a second argument which is the suffix (typically a domain name) to be added to all usernames when attempting to bind. All other children settings within this directive are not required, but may be needed depending on your LDAP configuration.

The LDAP URL should look like the following:

```
ldaps://ldap.example.org:1234/dc=example,dc=org
```

It begins with the protocol to be used when contacting the LDAP server -- either `ldap://` (for unencrypted LDAP) or `ldaps://` (LDAP over an encrypted SSL tunnel). We do not currently support StartTLS. If your LDAP server supports `ldaps://`, this option is more secure and ensures the username and passwords aren't being transmitted between Shiny Server and your LDAP server in plain text. Next, the hostname or IP of the LDAP server should follow the protocol. Bear in mind that the given hostname or IP *must* match the hostname associated with the SSL certificate if using `ldaps://`. If using a non-standard port (anything other than 389 for `ldap://` or 636 for `ldaps://`), you can follow the hostname with a colon and the port number which should be used. Following the hostname/port should be a forward slash then the root DIT of the directory to use.

There are a variety of settings which will be used to guide Shiny Server Pro's interaction with your LDAP server. These will all be defined inside of an `auth_ldap` or `auth_active_dir` setting and are described below. Be aware that the primary difference between `auth_ldap` and `auth_active_dir` is the default values assigned to these settings. If any of these settings are not specified, the default values will be used.

Some of the settings may include dynamic values which will be listed under the "Variables Available" section in each. These variables are enclosed in curly braces ("`{`" and "`}`") and will be replaced with their value before being used.

4.6.1 `base_bind` - Pro Only!

By default, the `auth_ldap` and `auth_active_dir` directives instruct Shiny Server Pro to use single-bind LDAP authentication for username and password validation. [Double-bind LDAP authentication](#) is used when the `base_bind` directive is defined. This directive specifies a user DN and password for the initial LDAP bind operation. The authenticated connection allows Shiny Server Pro to search for a DN associated with the user attempting login. The discovered DN is subsequently provided to a second bind operation.

```
base_bind <user-dn> <password-or-filepath>;
```


The `user_filter` directive is used during the user DN search and must be defined whenever using `base_bind` with `auth_ldap`. The `auth_active_dir` directive provides a default value to `user_filter` but `auth_ldap` does not.

Variables Available

- `{root}` -- The root DIT provided to the parent `auth_ldap` or `auth_active_dir` setting.

Here is an example `auth_ldap` specification using a searcher account during double-bind LDAP authentication with the required `user_filter` directive.

```
auth_ldap ldap://ldap.example.org/dc=example,dc=org {
  base_bind "uid=searcher,ou=People,{root}" "password-for-searcher";
  user_filter "uid={username}";
}
```

4.6.2 Securing the LDAP password - Pro Only!

In our earlier example, we specified the password as plain text in the Shiny Server Pro configuration file. We could alternatively use a fully-qualified file path containing only the LDAP password as in the following example:

```
auth_ldap ldap://ldap.example.org/dc=example,dc=org {
  base_bind "uid=searcher,ou=People,{root}" "/etc/shiny-server/ldap-base-bind-password.txt";
  user_filter "uid={username}"
}
```

In this case, Shiny Server Pro reads the LDAP password from the file `/etc/shiny-server/ldap-base-bind-password.txt`. Shiny Server Pro will retain root privileges when using a file to contain the LDAP password. This allows you to tighten the permissions around the file containing the LDAP password to only allow the root user to read it.

For example, we could secure the above file using:

```
# Change the file to be owned by the root user. {.pro}
chown root:root /etc/shiny-server/ldap-base-bind-password.txt

# Set the permissions so that only root can read the file. {.pro}
chmod 600 /etc/shiny-server/ldap-base-bind-password.txt
```

This prevents the LDAP password from Shiny applications or other users who are allowed to login to the server. If you expect that many users or untrusted users may be logging into or deploying apps on the server on which Shiny Server Pro is installed, we strongly recommend that you use this configuration to ensure the privacy of your LDAP password.

4.6.3 user_bind_template - Pro Only!

Shiny Server Pro will transform the username entered at the login screen before attempting to "bind" (authenticate) to the LDAP server using this template. In general (notable exception being Active Directory), this pattern should map the given username to the user's DN in the LDAP database. In systems where that may not be possible (such as Active Directory), you can use the `[user_filter]` setting to lookup the user's DN given their username. In either case, this setting should manipulate the given username into the username used to perform the LDAP bind operation.

Variables Available

- `{username}` -- The username entered by the user at the login screen.
- `{root}` -- The root DIT provided to the parent `auth_ldap` or `auth_active_dir` setting.

Default Value

- For `auth_ldap` -- `uid={username},ou=People,{root}`
- For `auth_active_dir` -- `{username}@example.org` (where "example.org" is the domain name you provided as the second argument to `auth_active_dir`).

4.6.4 group_search_base - Pro Only!

Defines the subtree in which groups are stored. This will be used as the root of all LDAP queries which attempt to find the groups of which a user is a member. The given value will be followed by a comma and the root DIT given in the `auth_ldap` or

auth_active_dir parent setting. If configured to use an empty string as the base, then the unmodified root DIT will be used as the group search base.

Default Value

- For auth_ldap -- ou=Groups
- For auth_active_dir -- cn=Users

4.6.5 group_filter - Pro Only!

The LDAP query to use in determining a user's group membership. The query should return all groups of which the given user is a member

The default for auth_ldap checks the uniqueMember field -- a setting from the groupOfUniqueMembers LDAP class. If you're using the posixGroup LDAP class to store your group memberships, you should likely set this value to memberUid={username}.

Variables Available

- {username} -- The username entered by the user at the login screen.
- {userDN} -- The computed DN for the entered user. If user_filter is unspecified, this will just be the value calculated using the pattern from user_bind_template. If user_filter is provided, then this will be the DN of the first object which matched the query specified in user_filter.
- {root} -- The root DIT provided to the parent auth_ldap or auth_active_dir setting.

Default Value

- For auth_ldap -- uniqueMember={username}
- For auth_active_dir -- member:1.2.840.1.13556.1.4.1941:={userDN}

4.6.6 group_name_attribute - Pro Only!

The attribute of the LDAP group object in which the group name is stored. This field will be compared to the entries in the required_group setting to determine whether or not a user should be granted access to a Shiny application.

Default Value

cn

4.6.7 trusted_ca - Pro Only!

By default, Shiny Server Pro trusts many standard SSL Certificate Authorities (CAs) (such as Verisign). If your organization uses a non-trusted Certificate Authority to sign its SSL certificates, you will need to explicitly tell Shiny Server Pro to trust this CA's certificate. You can do this by placing the CA's certificate (in PEM format) in a file on your machine and pointing this setting to that file. You can add multiple trusted CAs (space-delimited) if you desire.

If this value is provided, the standard list of trusted CAs will be overridden with the provided certificate.

Default Value

None.

4.6.8 check_ssl_ca - Pro Only!

When using LDAP over SSL (ldaps://), Shiny Server Pro will check that the SSL certificate was signed by a recognized Certificate Authority (CA) (though this can be modified using the trusted_ca setting). check_ssl_ca can be used to disable the checking of CAs entirely. By default, this option will be 'true' to enable SSL CA validation.

Disabling these checks is necessary if your LDAP server uses a self-signed certificate, for instance. Otherwise, this setting should typically be left enabled as disabling it will largely defeat the security gained by using SSL.

Default Value

true

4.6.9 user_filter - Pro Only!

Some systems (notably many Active Directory implementations) do not use the username as a part of the user's DN. In such systems, it may be necessary to perform an extra LDAP query after binding to determine the user's DN based on their username before group membership can be determined. This setting stores the LDAP filter used to find the user object which matches the entered username.

Using the default provided for auth_active_dir (userPrincipalName={userBind}), as an example. Shiny Server Pro will attempt to

bind to the LDAP server using the given username (after being manipulated as defined in `user_bind_template`) and password. If successful, it will then search for an object whose attribute `userPrincipalName` matches the username manipulated by `user_bind_template`. If found, the returned object's DN will be made available to the `group_filter` as the `{userDN}` variable.

Default Value

- For `auth_ldap` -- N/A
- For `auth_active_dir` -- `userPrincipalName={userBind}`

4.6.10 `user_search_base` - Pro Only!

Defines the subtree in which users are stored. This will only be used if `user_filter` is not empty, as that is the only time that Shiny Server Pro would need to search for a user. If provided, this will be used as the root of the LDAP queries used to lookup a user by his or her username. The given value will be followed by a comma and the root DIT given in the `auth_ldap` or `auth_active_dir` parent setting. If configured to use an empty string as the base, then the unmodified root DIT will be used as the user search base.

Default Value

- For `auth_ldap` -- `ou=People`
- For `auth_active_dir` -- `cn=Users`

4.7 Required Users & Groups - Pro Only!

4.7.1 User Authentication - Pro Only!

The `required_user` setting controls which users should be able to access a particular server, location, or admin, as in:

```
location /app1 {  
    required_user kim tom;  
}
```

This directive would enable authentication on the specified location, which would present any new, un-authenticated visitor to the application with a login page. If the visitor entered valid credentials against your configured authentication mechanism, then he or she would be checked against the list of allowed users for this application. If found (i.e. if the user name was either kim or tom), the visitor would be granted access to the application. Otherwise, he or she would be unable to access the application.

There are two special cases for `required_user`. Including the user named `*` will admit any user who is able to login. Additionally, including the user `**` will admit any visitor, whether or not they are logged in; this is the default behavior if no `required_user` setting is specified.

It is encouraged that any server accepting user credentials be secured via [SSL](#).

4.7.2 Group Authentication - Pro Only!

The `required_group` setting allows the administrator to control access based on the groups of which users are a member. In this way, entire sets of users can be granted access to a particular location in one line in the configuration file.

Flat-file authentication does not support the notion of user groups, but this feature can be used with LDAP or Active Directory by adding a restriction like the following:

```
location /app1 {  
    required_group shinyUsers admins;  
}
```

This configuration would grant access to any user who is a member of either the `shinyUsers` or `admins` group.

4.8 SSL - Pro Only!

In Shiny Server Professional, SSL encryption is available for any configured server and in the admin interface. SSL is a means of encrypting traffic between web clients and servers and should be used for any application that will accept or transmit sensitive information such as passwords. Untrusted SSL certificates can be generated locally by anyone for free; signed SSL certificates can be obtained from authorized certificate authorities fairly inexpensively. Untrusted SSL certificates will still encrypt the traffic

between a client and server, but (in most browsers), the client will need to proceed through a warning about the untrusted nature of the SSL certificate in use. SSL can be enabled on a server by configuring the two required parameters: the path to the SSL key, and the path to the SSL certificate.

```
server {  
  # Instruct this server to listen on port 443, the default port for HTTPS  
  # traffic  
  listen 443;  
  ssl /etc/shiny-server/ssl-key.pem /etc/shiny-server/ssl.cert;  
  
  ...  
}
```

If a valid SSL key and certificate were provided, a user would now need to preface their request with "https://" to be able to visit any application hosted within this server or on the admin. It is highly recommended that any Shiny Server Pro configuration which expects to accept its users' LDAP or Active Directory credentials use SSL to secure the usernames and passwords on their applications and admin interface.

4.9 Proxied Headers - Pro Only!

Typically, HTTP headers sent to Shiny Server will not be forwarded to the underlying Shiny application. However, Shiny Server Professional is able to forward specified headers into the Shiny application using the `whitelist_headers` configuration directive, which can be set globally or for a particular server or location.

```
server {  
  listen 3838;  
  whitelist_headers myorg_username myorg_groupnames myorg_privileges;  
  
  ...  
}
```

The above configuration would allow the three listed `myorg_` headers (case-insensitive) to be forwarded from the originating HTTP traffic into the Shiny session. In Shiny, these headers would be converted to capital letters and prefaced with `HTTP_` to designate that the headers originated in an HTTP request. They would then be available as objects in the `session$request` environment inside an application's `server.R` file. For example, the `myorg_username` object could be accessed via the following `server.R` file.

```
library(shiny)  
shinyServer(function(input, output, session){  
  output$username <- renderText({  
    session$request$HTTP_MYORG_USERNAME  
  })  
})
```

Organizations which employ a centralized authentication system not supported by Shiny Server may use this feature to leverage a proxied authentication system. In such a model, traffic destined for Shiny Server Professional would first be sent through an authenticating proxy which would take care of ensuring that the user is authenticated and the appropriate headers are set indicating the user, groups, or permissions. The application deployed on Shiny Server Professional would then use these headers to determine whether the user should be allowed access to the application, or a subset of the features of the application.

Be aware that, in such a model, the entirety of the `ui.R` file will be run and passed to the client before the server-side has a chance to make any decision about what privileges the current user should have given their HTTP headers (if any). So the proxy must be responsible for authenticating or rejecting visitors prior to forwarding any traffic to Shiny Server Professional if any sensitive information would be disclosed by a user merely opening the UI. Additionally, precautions should be taken to ensure that the Shiny Server machine is not accessible directly by users, but is instead *only* accessible through the authenticating proxy. Otherwise, an attacker would merely need to provide fraudulent HTTP headers when accessing Shiny Server in order to imitate another user or claim particular privileges.

5 Monitoring the Server - Pro Only!

5.1 Admin - Pro Only!

5.1.1 Configuration - Pro Only!

The Admin interface provided as a part of Shiny Server Professional is disabled by default. To enable it, use the admin setting and specify the port number on which the admin interface should run.

```
admin 3957 {  
  required_user tom sarah;  
}
```

The above would enable the admin interface on port 3957. This interface would then be available from any modern web browser on that port. First-time visitors would be presented with a login page. If the visitor were able to authenticate themselves to the Shiny Server application as either tom or sarah, they would be granted access to the admin dashboard.

The required user must be one connected to Shiny Server as described in [Authentication & Security](#).

You can also leverage [SSL](#) to secure the usernames and passwords accepted on the admin login screen.

5.1.2 Organization - Pro Only!

The admin interface offers a real-time view into the current state of your Shiny Server. There is no need to refresh the page periodically, as the information is updated as it happens. The interface is divided into four main sections.

Dashboard

Initially when logging in, the user will be presented with the Dashboard view. This view shows a basic overview of the state of Shiny Server, including the current RAM and CPU utilization, the number of connections, and historical RAM and CPU saturation.

Applications

The highest-level grouping in the admin interface is the notion of an "Application." An application represents a Shiny App stored in a particular directory on your server. There may be multiple R processes behind one application due to application restarts, but all of these R processes are considered part of the same application. Initially, you will be presented with a table depicting an overview of the various directories on your server in which applications are stored, the current number of processes associated with that application, and the number of open connections to that application.

Shiny Server does not crawl the disk to discover all the potential Shiny applications which could be spawned. Instead, it only becomes aware of the existence of a Shiny application once its visited.

Clicking on any "App Directory" in the table will take you to the detailed view for that application. On this page, you'll be able to track current and historical RAM usage, along with CPU utilization, active processes, and open connections.

You'll also see the current and historical "latency" for this application. Keep in mind that, because R is single-threaded, it is only able to compute one request at a time. Thus, while one user's request is being computed, any other requests that are assigned to this R process will be inserted into a queue until the R process becomes free. Only at that point, will computation begin on the second request. The latency of an application represents the delay between when a request initially arrived at Shiny Server and the moment the assigned R process began computation to fulfill that request. A Shiny application that is not computationally intensive and doesn't have many active users may see a latency of just a few milliseconds. Conversely, an application which is computationally intensive and/or overwhelmed with more users than the associated process(es) can handle will see a much higher latency.

Processes

The "Processes" tab provides access to information about the underlying R Shiny processes which support the various applications being hosted on this server. The overview table reveals with which application each process is associated, the latency of this application, and its current memory consumption. Clicking on the ID of the process will take you to the detailed view for this process. Here you can see the current CPU utilization, open connections, and various other information used to help understand this process.

Connections

This tab displays all of the connections currently open between Shiny Server and its visitors. You can see which application the connection is accessing, when it was opened, its current status (whether it is idle or awaiting the results of some analysis), the protocol being used, whether or not the connection is encrypted, the IP address associated with the connection, and -- if relevant -- the username to whom the connection belongs. You can also filter to only those connections associated with a particular app directory or status.

The protocol each connection is using will be displayed in a blue oval at the far left edge of the table using some abbreviation. If

you hover the cursor over the oval, you'll see the full protocol name. You can find more details on the various protocols supported by Shiny Server in the section on [Specifying Protocols](#).

5.1.3 Killing Processes and Connections - Pro Only!

The Admin interface also provides the opportunity to kill individual processes or connections. Alongside any table in which a process is listed, if you point your mouse at an application or connection, you will see a button appear on the right side which allows you to kill that app/connection. Clicking this button will present you with a confirmation dialogue box which asks you to confirm that you do desire to disconnect this connection or kill this process. If you confirm, that process or connection will be terminated.

If disconnecting a connection, the user will suddenly be disconnected from the server and the application will appear "greyed out." Keep in mind that, in some Shiny applications, this could forfeit some of the work being done by this user on that application. If disconnecting a process, any users with connections open to that process will be suddenly disconnected. Again, this potentially could cause the user to lose some data or work if they were using a complex Shiny application.

5.2 Graphite - Pro Only!

The [Admin](#) interface is often the preferred way to monitor the performance of a single Shiny Server Pro server. However, for organizations that manage a multitude of servers or multiple Shiny Server Pro instances, it may be convenient to send server monitoring data to a centralized database, rather than local RRD files. For such organizations, Shiny Server Pro can output its metrics to [Graphite](#) (or any other engine compatible with the Carbon protocol). This can be enabled in addition to, or in place of, local RRD file monitoring. Note that if RRD monitoring is disabled, no data will be recorded for later inspection in the Admin interface.

To enable Graphite metrics, use the `graphite_enabled` configuration directive. This setting accepts two optional arguments: the address of the Graphite server (host) and the port on which Carbon -- the component of Graphite responsible for the intake of new data -- is listening on that server. By default, these will be set to 127.0.0.1 and 2003, respectively.

RRD metrics are enabled by default and used by the Admin interface. If you choose to disable RRD metrics, you can use the `rrd_disabled` setting. If this setting is present in your configuration file, RRD metrics will not be recorded; in such a configuration, the admin interface cannot be enabled.

Changes to either of these settings will not take effect until the server is restarted.

5.3 Health Check Endpoint - Pro Only!

Some advanced deployment scenarios are able to leverage a "health check endpoint" which provides an automated way to ensure that Shiny Server is online and responsive.

The `/ping` endpoint has been removed from Shiny Server as of version 1.1.0 in favor of the `/__health-check__` endpoint.

The URL `/__health-check__` (note the double underscores on either side) can be used on any port Shiny Server is instructed to listen on in order to test the health of Shiny Server. If the server is online and responsive, it will respond an HTTP 200 code and various information about the server including the number of current connections, the current CPU load, and other information about the server. By default, the page will look something like this:

```
server-version: 1.2.3.4
active-connections: 8
active-apps: 2
active-processes: 3
cpu-percent: 13
memory-percent: 49
swap-percent: 39.1
load-average: 1.01953125
```

5.3.1 Customizing Responses - Pro Only!

Shiny Server Pro uses a templating system to determine how to respond to visits to the health-check URL. By default, it uses the following template:

```
server-version: #server-version#
active-connections: #active-connections#
```

```
active-apps: #active-apps#
active-processes: #active-processes#
cpu-percent: #cpu-percent#
memory-percent: #memory-percent#
swap-percent: #swap-percent#
load-average: #load-average#
```

As you might infer, the values encompassed by # symbols are treated as variables that will be substituted for their true values before being returned. The template above shows all of the variables that are currently available to the endpoint.

You can specify whatever format you like for these variables by creating a custom template in /etc/shiny-server/health-check. This file will be used to structure the response to the health check after substituting the available variables. For instance, you could configure the endpoint to return this data in XML by placing the following content in the file:

```
<?xml version="1.0" encoding="UTF-8"?>
<health-check>
  <server-version>#server-version#</server-version>
  <active-connections>#active-connections#</active-connections>
  <active-apps>#active-apps#</active-apps>
  <active-processes>#active-processes#</active-processes>
  <cpu-percent>#cpu-percent#</cpu-percent>
  <memory-percent>#memory-percent#</memory-percent>
  <swap-percent>#swap-percent#</swap-percent>
  <load-average>#load-average#</load-average>
</health-check>
```

5.3.2 Providing Multiple Health-Check Endpoints - Pro Only!

The file at /etc/shiny-server/health-check specifies the default template to use for the health check. Shiny Server Pro supports multiple health check endpoints which can be accessed using different URLs. For instance, you can setup a new endpoint that responds using JSON using the following template:

```
{
  'server-version' : '#server-version#',
  'active-connections': #active-connections#,
  'active-apps': #active-apps#,
  'active-processes': #active-processes#,
  'cpu-percent': #cpu-percent#,
  'memory-percent': #memory-percent#,
  'swap-percent': #swap-percent#,
  'load-average': #load-average#
}
```

Store this text in a new file at /etc/shiny-server/health-check-json (the file is suffixed by a *hyphen* then some custom string). This endpoint would then be available at `http://<server-address-and-port>/__health-check__<custom-string>.json` (the URL is suffixed by a *forward slash* then the same custom string).

6 Licensing & Activation - Pro Only!

6.1 Product Activation - Pro Only!

6.1.1 Activation Basics - Pro Only!

When Shiny Server Professional is first installed on a system it operates in evaluation mode for a period of time and then subsequently requires activation for continued use. To determine the current license status of your system you can use the following command:

```
$ sudo /opt/shiny-server/bin/license-manager status
```

After purchasing a license to Shiny Server Professional you'll receive a product key that can be used to activate the license on a

given system. You can perform the activation as follows:

```
$ sudo /opt/shiny-server/bin/license-manager activate <product-key>
$ sudo restart shiny-server # or "sudo /sbin/service shiny-server restart" or "sudo systemctl restart shiny-server"
```

Note that you need to reload or restart the server for licensing changes to take effect.

If you want to move your license of Shiny Server Professional to another system you should first deactivate it on the system you are moving from. For example:

```
$ sudo /opt/shiny-server/bin/license-manager deactivate
```

6.2 Connectivity Requirements - Pro Only!

In order to activate or deactivate Shiny Server Professional internet connectivity is required for communication with the licensing server. If your server is behind an internet proxy or not connected to the internet at all this section describes what's required to successfully activate.

6.2.1 Proxy Servers - Pro Only!

If your server is behind an internet proxy you may need to add an additional command line flag indicating the address and credentials required to communicate through the proxy. Note however that this may not be necessary if either the `http_proxy` or `all_proxy` environment variable is defined (these are read and used by Shiny Server Professional when available).

If you do need to specify a proxy server explicitly you can do so using the `--proxy` command line parameter. For example:

```
$ sudo /opt/shiny-server/bin/license-manager --proxy=http://127.0.0.1/ activate <product-key>
```

Proxy settings can include a host-name, port, and username/password if necessary. The following are all valid proxy configurations:

```
http://127.0.0.1/
http://127.0.0.1:8080/
http://user:pass@127.0.0.1:8080/
```

If the port is not specified, the license manager will default to using port 1080.

6.2.2 Offline Activation - Pro Only!

If your system has no connection to the internet it's also possible to perform an offline activation. To do this, you first generate an offline activation request as follows:

```
$ sudo /opt/shiny-server/bin/license-manager activate-offline-request <product-key>
```

Executing this command will print an offline activation request to the terminal which you should copy and paste and then send to RStudio customer support (support@rstudio.com). You will receive a reply with a file attachment that can be used to activate offline as follows:

```
$ sudo /opt/shiny-server/bin/license-manager activate-offline <activation-file>
$ sudo reload shiny-server # or "sudo /sbin/service shiny-server reload"
```

Note that you need to reload or restart the server in order for licensing changes to take effect.

If you want to move your license of Shiny Server Professional to another system you can also perform license deactivation offline. You can do this as follows:

```
$ sudo /opt/shiny-server/bin/license-manager deactivate-offline
```


Executing this command will print an offline deactivation request to the terminal which you should copy and paste and then send to RStudio customer support (support@rstudio.com).

You can also perform an offline check of your current license status using the following command:

```
$ sudo /opt/shiny-server/bin/license-manager status-offline
```

7 Appendix

7.1 Quick Start

We recommend reading through the relevant sections of this guide to gain a complete understanding of how to operate Shiny Server. If you're just looking to get running quickly, you can follow one of the guides below.

All of the guides will help you customize your Shiny Server configuration file which will be stored at `/etc/shiny-server/shiny-server.conf`. By altering this file, you can control exactly how Shiny Server runs.

The following guides are available:

- [Host a directory of applications](#) - This is the default configuration that Shiny Server will use until you provide a custom configuration in `/etc/shiny-server/shiny-server.conf`. This guide will show you how to serve multiple applications from a single directory on disk -- `/srv/shiny-server/`. If you're not sure where to start, we recommend this guide.
- [Let users manage their own applications](#) - This guide will serve Shiny applications created by your users and stored in their home directories.
- [Run Shiny Server on multiple ports](#) - This guide will introduce a multi-server configuration which demonstrates how to run two distinct services on different ports on your server. It also configures one server with multiple locations.

PRO

And for Pro features:

- [Require user authentication on an application](#) - This example will help you setup a local username/password database and use it to secure your applications.
- [Host a secure Shiny Server](#) - If you're looking to setup a secure instance of Shiny Server which is encrypted with HTTPS (SSL/TLS), this guide will show you how.
- [Host an Application Supported by Multiple R Processes](#) - This guide will demonstrate how to scale a Shiny application to multiple R processes.

7.1.1 Host a directory of applications

Shiny Server is configured by a file stored at `/etc/shiny-server/shiny-server.conf`. To more conveniently use this Quick Start guide, we provide a convenient executable in `/opt/shiny-server/bin/deploy-example` that can copy the example configuration file associated with this Quick Start to that location and install a directory of example applications. Be sure to backup your current configuration file if you've done any customizing, then run the following command:

```
sudo /opt/shiny-server/bin/deploy-example default
```

This configuration expects your Shiny applications to be hosted in `/srv/shiny-server/`. The installer will also place one sample application in `/srv/shiny-server/sample-apps/hello/`.

By default, Shiny Server listens on port 3838, so the default application will be available at `http://<server-address>:3838/sample-apps/hello/`.

If you wanted to change the port or the directory Shiny Server uses, you would adjust the configuration file. To begin customizing your server, open `/etc/shiny-server/shiny-server.conf` in your preferred text editor and modify the line which says listen 3838. This is the line that instructs your server to listen to port 3838 -- change it to another port such as 12345 and save the file. You'll need to restart Shiny Server to make it aware of the new configuration file. The [Stopping and Starting](#) section explains this in more detail. For now, you'll run this command if you're on a distribution which supports Upstart (most recent distributions, including Ubuntu 12 and later and CentOS 6):

```
sudo restart shiny-server
```

If you're on an older system that relies on `init.d` (such as CentOS 5), you'll use the following command instead:

```
sudo /sbin/service shiny-server restart
```


You should now find that Shiny Server is no longer running on port 3838, but is on whichever new port you selected.

7.1.2 Let users manage their own applications

Some administrators of Shiny Server would prefer to give users -- or some trusted subgroup of users -- the ability to manage and host their own Shiny applications. This allows users to work directly with their applications rather than requiring an administrator to deploy and update applications on their behalf. This guide will document how to create such an environment using the `user_dirs` directive documented in the section named [Host Per-User Application Directories](#).

To more conveniently use this Quick Start guide, we provide a convenient executable in `/opt/shiny-server/bin/deploy-example` that can configure Shiny Server to use the example configuration file associated with this guide (by placing it in `/etc/shiny-server/shiny-server.conf`). It will also install a directory of example applications. Be sure to backup your current configuration file if you've done any customizing, then run the following command:

```
sudo /opt/shiny-server/bin/deploy-example user-dirs
```

This configuration enables all users on the system to host their own applications by creating a `ShinyApps` directory in their home directory; you can always alter the configuration file later to tailor it to meet your needs. You can deploy an application under your username by copying an example that comes with Shiny Server into your own `ShinyApps` directory. To do this, copy that entire `/opt/shiny-server/samples/sample-apps/hello/` directory into `~/ShinyApps` using the following command:

```
mkdir ~/ShinyApps
sudo cp -R /opt/shiny-server/samples/sample-apps/hello ~/ShinyApps/
```

If that command succeeds, you've successfully installed a `user_dirs` Shiny application!

By default, Shiny Server listens on port 3838, so your new application will be available at `http://<server-address>:3838/<your_username>/hello` where `<your_username>` is your Linux username.

If you wanted to restrict the privilege of running Shiny applications to a particular set of users on your system, you can uncomment the `members_of` line in the configuration file. You can then specify the Linux group name of which a user must be a member before he or she would be allowed to host Shiny applications.

7.1.3 Run Shiny Server on multiple ports

This guide will show how to run Shiny Server on multiple ports simultaneously and will also show how to run a single server with multiple locations.

To more conveniently use this Quick Start guide, we provide a convenient executable in `/opt/shiny-server/bin/deploy-example` that can configure Shiny Server to use the example configuration file associated with this guide (by placing it in `/etc/shiny-server/shiny-server.conf`). It will also install a directory of example applications. Be sure to backup your current configuration file if you've done any customizing, then run the following command:

```
sudo /opt/shiny-server/bin/deploy-example multi-server
```

This configuration will run two servers: one on port 3838, and one on port 4949. As you can see in the configuration file, a variety of different locations are defined with different hosting models.

The server running on port 3838 defines two locations. The first for the URL `/users` which makes user applications available using the `user_dirs` hosting model as documented in the section entitled [Host Per-User Application Directories](#). The second location is at the URL `/example1` and hosts a single application which should be stored in `/srv/shiny-server/sample-apps/hello/`. An application is provided here during installation.

However you'll want to copy that application to your home directory, as well, so you have something to view in the `/users` location. You can copy that application to your home directory with the following commands:

```
# Copy to your home directory's 'ShinyApps' folder
mkdir ~/ShinyApps
cp -R /srv/shiny-server/sample-apps/hello ~/ShinyApps/
```

You should now be able to access applications in a variety of ways:

- On the server running on port 3838, you should have a `user_dirs` location at `/users` hosting your personal applications. You

can access the one deployed earlier at a URL like `http://<server-address>:3838/users/<your_username>/hello` where `<your_username>` is your Linux username.

- Also on the server running on port 3838, you have a single application available at the `/example1` location. So the URL like `http://<server-address>:3838/example1` should point you to the same example.
- Finally, the server on port 4949 will host any application defined in `/srv/shiny-server`. You can access one that comes with Shiny Server at a URL like `http://<server-address>:4949/sample-apps/hello` which, coincidentally, should point you to the same application you just loaded a moment ago.

7.1.4 Require user authentication on an application - Pro Only!

Shiny Server Pro comes with multiple different methods of authenticating users (as described in the chapter on [Authentication & Security](#)). Out of the box, it's configured to use a [Flat-File Authentication](#) database stored in `/etc/shiny-server/passwd`. This database is initially empty; to create a user named `admin`, execute the following command:

```
$ sudo /opt/shiny-server/bin/sspasswd /etc/shiny-server/passwd admin
```

then enter the new password for the `admin` user when prompted. Once completed, you'll have a single user defined in your database named `admin`.

You should now be able to login to your administrative dashboard at a URL like `http://myserver.com:4151/` (after replacing `myserver.com` with your server's IP or hostname) using the username `admin` with the password you just created. See the chapter entitled [Admin](#) for more details on how to take advantage of this interface.

After defining usernames and passwords in this password database, you can also restrict access to servers and locations using the `required_user` directive. To more conveniently use this Quick Start guide, we provide a convenient executable in `/opt/shiny-server/bin/deploy-example` that can configure Shiny Server to use the example configuration file associated with this guide (by placing it in `/etc/shiny-server/shiny-server.conf`). It will also install a directory of example applications. Be sure to backup your current configuration file if you've done any customizing, then run the following command:

```
sudo /opt/shiny-server/bin/deploy-example auth
```

Because of the specifications provided in our current configuration file (in `/etc/shiny-server/shiny-server.conf`), the example Shiny application that comes with Shiny Server requires authentication -- you must be logged in as the user named `"admin"` to access it. If you already visited the Administrative Dashboard on port 4151 mentioned earlier, then you may already be logged in. But you can test your configuration by either logging out or using another browser/computer to confirm that you're unable to access the application you just copied hosted at a URL like `http://<server-address>:3838/sample-apps/hello` without logging in. However, once you login as the user named `"admin"` using the password you defined above, you will be granted access to this application.

Note that, because the configuration in use is structured like so:

```
server {
  listen 3838;

  location / {
    ...

    location /sample-apps/hello {
      required_user admin;
    }
  }
}
```

only the `hello` location requires authentication; all other applications are freely available. If you were to move the `required_user` directive to the parent location, however, all applications defined in that location would require being authenticated as the user named `"admin"`.

7.1.5 Host a secure Shiny Server - Pro Only!

Shiny Server Professional is able to serve Shiny applications using SSL/TLS -- an encrypted channel between your server and your clients.

The configuration we'll use expects your SSL key and certificate to be available in `/etc/shiny-server/server.key` and

/etc/shiny-server/server.cert, respectively. If you have an existing SSL certificate stored elsewhere, you can update the configuration file at /etc/shiny-server/shiny-server.conf to point to the appropriate files. If you don't have an SSL certificate available, you can setup a "self-signed" certificate. This certificate will allow you to encrypt traffic between your clients and server but will not be "trusted" by most browsers. Thus, when a user visits a page secured via a self-signed certificate, the user will get a warning advising them not to proceed. Nonetheless, we can demonstrate how to create a self-signed certificate and integrate it into Shiny Server. The details of creating and signing SSL certificates is outside the scope of this guide, but the commands are given below for your convenience.

In order to create a certificate, you'll need to have the openssl library installed on your server. You can then create an SSL key using the following command.

```
sudo openssl genrsa -out /etc/shiny-server/server.key 1024
```

You'll then generate a certificate signing request (CSR) using the following command:

```
sudo openssl req -new -key /etc/shiny-server/server.key \  
-out /etc/shiny-server/server.csr
```

(Where a command includes a \ you can either enter that literally as it appears in this document or you can omit it and enter the whole command on a single line.) This command will ask you a variety of information about your organization which you can provide or omit, if you just plan to sign this certificate yourself.

You can then sign that request yourself using this command:

```
sudo openssl x509 -req -days 3650 \  
-in /etc/shiny-server/server.csr \  
-signkey /etc/shiny-server/server.key \  
-out /etc/shiny-server/server.cert
```

To more conveniently use this Quick Start guide, we provide a convenient executable in /opt/shiny-server/bin/deploy-example that can configure Shiny Server to use the example configuration file associated with this guide (by placing it in /etc/shiny-server/shiny-server.conf). It will also install a directory of example applications. Be sure to backup your current configuration file if you've done any customizing, then run the following command:

```
sudo /opt/shiny-server/bin/deploy-example ssl
```

It is a good idea to inspect the log file at this point to ensure that there were no problems using your SSL key and certificate. The server log file is stored in /var/log/shiny-server.log. You can preview the last 20 lines of this file using the following command.

```
sudo tail -n 20 /var/log/shiny-server.log
```

You can ensure here that there were no errors in using your SSL configuration.

You should now be able to connect to your server from a browser by visiting a URL like `https://<server-address>:3939/sample-apps/hello`. Note that this guide uses the port 3939, unlike the other guides and that you *must* specify the `https://` protocol when visiting the page. If you signed the certificate yourself, your browser will likely prompt you about the untrusted certificate. If you instruct your browser to accept the certificate regardless, you'll be taken to your application which now is secured via SSL/TLS encryption.

7.1.6 Host an Application Supported by Multiple R Processes - Pro Only!

You can configure Shiny Server Professional to run multiple Shiny processes for a single application. You could evenly balance incoming traffic across up to three processes (the default), or even create a separate Shiny process for each incoming user. The details for configuring this setup is discussed in the section on the [Utilization Scheduler](#). To more conveniently use this Quick Start guide, we provide a convenient executable in /opt/shiny-server/bin/deploy-example that can configure Shiny Server to use the example configuration file associated with this guide (by placing it in /etc/shiny-server/shiny-server.conf). It will also install a directory of example applications. Be sure to backup your current configuration file if you've done any customizing, then run the following command:

```
sudo /opt/shiny-server/bin/deploy-example multi-proc
```

This configuration includes the following line to specify that a Utilization Scheduler should be used when serving Shiny applications.

```
utilization_scheduler 1 .7 5;
```

This configuration dictates that a Shiny process should only have one active Shiny session and that there should be at most 5 Shiny processes per application. This configuration would be appropriate for an application that were computationally intensive. In such an application, one session may tie up the Shiny process for minutes at a time while doing computation which would create an unpleasant experience for other users who try to connect during that window. In this way, we can ensure that each user gets his or her own designated Shiny process.

You should be able to connect to the sample application that Shiny Server installs on your server from a browser by visiting a URL like `http://<server-address>:3838/sample-apps/hello`.

In order to inspect what processes are being created and how they're performing, it is useful to open Shiny Server's Admin Dashboard. Before you can login to the dashboard, you'll need to create a user who is allowed to access the Dashboard. Shiny Server Pro comes with multiple different methods of authenticating users (as described in the chapter on [Authentication & Security](#)). Out of the box, it's configured to use a [Flat-File Authentication](#) database stored in `/etc/shiny-server/passwd`. This database is initially empty; to create a user named admin, execute the following command:

```
$ sudo /opt/shiny-server/bin/sspasswd /etc/shiny-server/passwd admin
```

then enter the new password for the admin user when prompted. Once completed, you'll have a single user defined in your database named admin.

You should now be able to login to your administrative dashboard at a URL like `http://<server-address>:4151/` using the username admin with the password you just created. See the chapter entitled [Admin](#) for more details on how to take advantage of this interface.

You will be able to monitor the processes that are being created in real-time using the Dashboard. You can also drill down into each process to verify that the connections are being assigned to their own private Shiny processes as we expected.

7.2 Configuration Settings

Listed below are all the directives that are supported in Shiny Server config files.

Applies to indicates the kind of parent scope that this directive normally appears inside.

Inheritable means that you can put this directive at a higher level in the hierarchy and it will be inherited by any children to which it might apply. Inherited directives can be overridden by using the directive again in a child scope.

run_as

The user the app should be run as. This user should have the minimal amount of privileges necessary to successfully run the application (i.e. read-only access to the Shiny application directory). Note that this directive cannot be used with [user_apps](#), as [user_apps](#) always run as the user who owns the application.

Parameter	Data type	Type	Description	Default
users	String	multiple	The username that should be used to run the app. If using user_dirs , this can also be a special keyword such as <code>:HOME_USER:</code> which will instruct <code>user_dirs</code> to run as the user in whose home directory the application exists. Alternatively it could be <code>:AUTH_USER:</code> which will run the application as the user who the visitor is logged in as. If using a special keyword like <code>:HOME_USER:</code> or <code>:AUTH_USER:</code> , you can specify an additional username afterwards which will be used when the special cases are not available, such as if the user is not logged in and the location is not served by <code>user_dirs</code> .	

Applies to: Top-level, [server](#), [location](#)

Inheritable: Yes

access_log

The file path of the HTTP access log.

--	--	--	--	--

Parameter	Data type	Type	Description	Default
path	String	required	The file path where the access log should be written	
format	String	optional	The log file format; see the Connect documentation (http://www.senchalabs.org/connect/logger.html) under "Formats"	default

Applies to: Top-level

Inheritable: Yes

server

Declares an HTTP server. You need one of these for each port/IP address combination this Shiny Server instance should listen on.

This directive has no parameters.

Applies to: Top-level

Inheritable: Yes

Child directives: [listen](#), [server_name](#)

listen

Directs the enclosing server scope to listen on the specified port/IP address combination.

Parameter	Data type	Type	Description	Default
port	Integer	required	Port to listen on	
host	String	optional	IPv4 address to listen on (* or 0.0.0.0 for all); hostnames are not currently allowed, please use raw IPv4 address only	*

Applies to: [server](#)

Inheritable: Yes

server_name

Directs the enclosing server scope to only honor requests that have the given host headers (i.e. virtual hosts).

Parameter	Data type	Type	Description	Default
names	String	multiple	The virtual hostname(s) to bind this server to	

Applies to: [server](#)

Inheritable: Yes

location

Creates a scope that configures the given URL as a website ([site_dir](#)), specific application ([app_dir](#)), autouser root ([user_apps](#)), autouser root with [run_as](#) support ([user_dirs](#)), or redirect to a different URL ([redirect](#)).

Parameter	Data type	Type	Description	Default
path	String	required	The request path that this location should match	

Applies to: [server](#), [location](#)

Inheritable: Yes

Child directives: [run_as](#), [location](#), [site_dir](#), [directory_index](#), [user_apps](#), [user_dirs](#), [app_dir](#), [redirect](#), [log_dir](#), [members_of](#), [application](#), [pam_sessions_profile](#), [exec_supervisor](#), [r_path](#)

site_dir

Configures the enclosing location scope to be a website that can contain both Shiny applications and unrelated static assets in a single directory tree.

Parameter	Data type	Type	Description	Default
rootPath	String	required	The path to the root directory of the website	

Applies to: [application](#), [location](#)
Inheritable: Yes

directory_index

When enabled, if a directory is requested by the client and an index.html file is not present, a list of the directory contents is created automatically and returned to the client. If this directive is not present in a custom config file, the default behavior is to disable directory indexes. However, it is enabled if no config file is present at all (in other words, the default config file has it enabled).

Parameter	Data type	Type	Description	Default
enabled	Boolean	required	Whether directory contents should automatically displayed	

Applies to: Top-level, [server](#), [location](#)
Inheritable: Yes

user_apps

DEPRECATED! This directive has been deprecated in favor of [user_dirs](#), which offers more flexibility with regards to the [run_as](#) configuration. Configures the enclosing location scope to be an autouser root, meaning that applications will be served up from users' ~/ShinyApps directories and all Shiny processes will run as the user in whose directory the application is found.

Parameter	Data type	Type	Description	Default
enabled	Boolean	optional	Whether this location should serve up all users' ~/ShinyApps directories	true

Applies to: [location](#)
Inheritable: Yes

user_dirs

Configures the enclosing location scope to be an autouser root, meaning that applications will be served up from users' ~/ShinyApps directories. This directive does respect an affiliated run_as setting, meaning that the applications will be executed as whichever user is configured in the applicable run_as setting. Note that many distributions, by default, will prohibit users from being able to access each other's home directories.

Parameter	Data type	Type	Description	Default
enabled	Boolean	optional	Whether this location should serve up all users' ~/ShinyApps directories	true

Applies to: [location](#)
Inheritable: Yes

app_dir

Configures the enclosing location scope to serve up the specified Shiny application.

Parameter	Data type	Type	Description	Default
path	String	required	The path to the Shiny application directory	

Applies to: [location](#)
Inheritable: Yes

redirect

Configures the enclosing location to redirect all requests to the specified URL.

Parameter	Data type	Type	Description	Default
url	String	required	The URL (or base URL) to redirect to	
statusCode	Integer	optional	The status code to send with the response (usually 301 for permanent redirects or 302 for temporary redirects)	302
exact	Boolean	optional	Whether to match on the URL exactly; if false, any subpaths will match as well	true

Applies to: [location](#)

Inheritable: Yes

log_dir

Directs the application to write error logs to the specified directory. Only applies to location scopes that are configured with [app_dir](#) or [site_dir](#), as [user_apps](#) (autouser) always writes error logs to ~/ShinyApps/log.

Parameter	Data type	Type	Description	Default
path	String	required	The path to which application log files should be written	

Applies to: Top-level, [server](#), [location](#)

Inheritable: Yes

members_of

Restricts a [user_apps](#) or [user_dirs](#) (autouser) scope to require membership in one or more groups (or, if no arguments are passed, lifts group restrictions from a [members_of](#) directive in a parent scope).

Parameter	Data type	Type	Description	Default
groups	String	multiple	Users must be a member of at least one of these groups in order to deploy applications; if no groups are provided, then all users are allowed	

Applies to: Top-level, [server](#), [location](#)

Inheritable: Yes

required_group

Restricts a location to require membership in one or more groups in order to access the application.

Parameter	Data type	Type	Description	Default
groups	String	multiple	Users must be a member of at least one of these groups in order to access the application	

Applies to: Top-level, [server](#), [location](#), [admin](#), [application](#)

Inheritable: Yes

required_user

Restricts a location to only be available to the specified users.

Parameter	Data type	Type	Description	Default
users	String	multiple	Users who are allowed to access the application	

Applies to: Top-level, [server](#), [location](#), [admin](#), [application](#)

Inheritable: Yes

auth_passwd_file

Configures the server to use the specified file as the username/password store. The file should contain zero or more lines in the format username + ":{script}" + hashed-password, where hashed-password is the scripted password in base64 format. Use the `spspasswd` utility to create and modify passwd files.

Parameter	Data type	Type	Description	Default
path	String	required	The file path where the username/password file can be found	

Applies to: Top-level

Inheritable: Yes

auth_duration

Approximate time in minutes that a user should be kept logged in, from the time that they finish using their last application. The default value is 120.

Parameter	Data type	Type	Description	Default
minutes	Float	required	The number of minutes; must be greater than or equal to 10 or applications may behave unpredictably	

Applies to: Top-level

Inheritable: Yes

google_analytics_id

Configure Google Analytics tracking code to be inserted in Shiny application pages.

Parameter	Data type	Type	Description	Default
gaid	String	required	The Google tracking ID, for example, UA-18988-1	

Applies to: Top-level, [server](#), [location](#), [application](#)

Inheritable: Yes

app_init_timeout

Defines the amount of time Shiny Server will wait for an R process to start before giving up.

Parameter	Data type	Type	Description	Default
timeout	Integer	required	The number of seconds to wait for the application to start.	

Applies to: Top-level, [server](#), [location](#), [application](#)

Inheritable: Yes

app_idle_timeout

Defines the amount of time an R process will persist with no connections before being terminated.

Parameter	Data type	Type	Description	Default
timeout	Integer	required	The number of seconds to keep an empty R process alive before killing it.	

Applies to: Top-level, [server](#), [location](#), [application](#)

Inheritable: Yes

app_session_timeout

Defines how long an idle user's session should remain connected to the server.

Parameter	Data type	Type	Description	Default
timeout	Integer	required	The number of seconds after which an idle session will be disconnected. If 0, sessions will never be automatically disconnected.	

Applies to: Top-level, [server](#), [location](#), [application](#)

Inheritable: Yes

simple_scheduler

A basic scheduler which will spawn one single-threaded R worker for each application. If no scheduler is specified, this is the default scheduler.

Parameter	Data type	Type	Description	Default
			The maximum number of requests to assign to this scheduler before it should start	

maxRequests	Integer	optional	returning rejecting incoming traffic using a '503 - Service Unavailable' message. Once this threshold is hit, users attempting to initialize a new session will receive 503 errors.	100
-------------	---------	----------	---	-----

Applies to: Top-level, [server](#), [location](#), [application](#)

Inheritable: Yes

utilization_scheduler

A multi-process scheduler which will load-balance incoming traffic to multiple R processes. Each incoming request will be routed to the running worker with the fewest open connections, unless the request is explicitly targetting a particular worker.

Parameter	Data type	Type	Description	Default
maxRequestsPerProc	Integer	required	The maximum number of simultaneous requests (web socket or HTTP) per R process.	
loadFactor	Float	required	A decimal value in [0, 1] defining the capacity at which a new R process should be pre-emptively spawned.	
maxProc	Integer	required	The maximum number of concurrent R processes to run for this application.	

Applies to: Top-level, [server](#), [location](#), [application](#)

Inheritable: Yes

ssl

The absolute paths to the files to use as the SSL key and certificate for this server.

Parameter	Data type	Type	Description	Default
key_path	String	required	The absolute path to the SSL key for this server.	
cert_path	String	required	The absolute path to the SSL certificate for this server.	

Applies to: [server](#), [admin](#)

Inheritable: Yes

allow_app_override

If present, will allow users to override the global defaults for a scheduler by customizing the parameters associated with a scheduler or even the type of scheduler used.

Parameter	Data type	Type	Description	Default
enabled	Boolean	optional	Whether or not this is enabled. Default is true.	true

Applies to: Top-level

Inheritable: Yes

admin

If present, will provide an administrative interface on the specified port allowing specified users to monitor deployed Shiny applications.

Parameter	Data type	Type	Description	Default
port	Integer	required	The port on which the admin interface will be made available.	

Applies to: Top-level

Inheritable: Yes

Child directives: [ssl](#), [template_dir](#)

whitelist_headers

Enumerates the HTTP headers that should be trusted and passed from the original HTTP request to the Shiny application.

Parameter	Data type	Type	Description	Default
-----------	-----------	------	-------------	---------

headers	String	multiple	The list of headers which should be trusted.	
---------	--------	----------	--	--

Applies to: Top-level, [server](#), [location](#), [application](#)

Inheritable: Yes

auth_google

Authenticate to a Google Application using a configured application ID and secret.

Parameter	Data type	Type	Description	Default
id	String	required	The Client ID available in your Google Apps Console.	
secret	String	required	The Client Secret available in your Google Apps Console, or the absolute path to the file storing the secret. If using a file path, Shiny Server will retain root privileges under the assumption that the file is only readable by the root user.	
filters	String	multiple	Used to limit the Google users who will be granted access to the system by filtering based on their email address. Filters may begin with a '+' or a '-' to indicate whether this is a 'positive' match (in which case these users should be granted access), or a 'negative' match (in which case any matching email should be rejected). These filters will iteratively be applied to the incoming email address until a match is found, at which point the user will either be admitted or rejected based on whether the filter was a 'positive' or 'negative' filter. A filter not beginning with a '+' or '-' is assumed to be a 'positive' filter. The asterisk can be used as a wildcard.	

Applies to: Top-level

Inheritable: Yes

auth_pam

Authenticate using the local server's PAM configuration.

Parameter	Data type	Type	Description	Default
forwardPassword	Boolean	optional	If true, will retain the user's password when they login so that it can be provided to PAM when spawning processes on this user's behalf. This can be used where PAM features such as Kerberos or pam_mount which require the user's password are in use.	

Applies to: Top-level

Inheritable: Yes

auth_ldap

Define how the LDAP (or Active Directory) server will be accessed.

Parameter	Data type	Type	Description	Default
url	String	required	The URL for the LDAP server and the root DIT of the directory. The URL can either use the unsecured 'ldap://' protocol, or the SSL-secured 'ldaps://' protocol, followed by the hostname or IP address of the LDAP server. Optionally, a port number can be provided after the LDAP server's address prefixed by a colon. If no port number is provided, the default port numbers of 389 and 636 for 'ldap://' and 'ldaps://' (respectively) will be used. This should be followed by a forward slash and the root DIT of the directory to use for authentication. For example, one configuration could be 'ldaps://ldap.example.org:1636/dc=example,dc=org'.	

Applies to: Top-level

Inheritable: Yes

auth_active_dir

Define how the Active Directory server will be accessed. This provides the same functionality as [auth_ldap](#), but tailors the defaults to be appropriate for Active Directory.

Parameter	Data	Type	Description	Default
-----------	------	------	-------------	---------

	type			
url	String	required	The URL for the Active Directory (AD) server and the root DIT of the directory. The URL can either use the unsecured 'ldap://' protocol, or the SSL-secured 'ldaps://' protocol, followed by the hostname or IP address of the AD server. Optionally, a port number can be provided after the LDAP server's address prefixed by a colon. If no port number is provided, the default port numbers of 389 and 636 for 'ldap://' and 'ldaps://' (respectively) will be used. This should be followed by a forward slash and the root DIT of the directory to use for authentication. For example, one configuration could be 'ldaps://ad.example.org:1636/dc=example,dc=org'.	
userSuffix	String	required	The suffix (typically the domain name) to be appended to usernames when attempting to bind to the AD server. The default user_bind_template will be set to '{{username}}@' followed by the userSuffix. This setting will be ignored if user_bind_template is explicitly set.	

Applies to: Top-level

Inheritable: Yes

Child directives: [ldap_timeout](#), [base_bind](#), [user_bind_template](#), [trusted_ca](#), [check_ssl_ca](#), [user_filter](#), [user_search_base](#), [group_search_base](#), [group_filter](#), [group_name_attribute](#)

ldap_timeout

Define the timeout for the parent LDAP connection. This timeout will be used when trying to connect to the LDAP server, and also when trying to query the server once connected. If not provided, a default timeout of 10 seconds will be used.

Parameter	Data type	Type	Description	Default
timeout	Integer	required	The number of seconds to wait for an LDAP query or connections before labeling it a failure. If 0, will never timeout.	

Applies to: [auth_ldap](#), [auth_active_dir](#)

Inheritable: Yes

base_bind

Define credentials used to perform the initial LDAP bind request for double-bind authentication. If not provided, single-bind authentication is performed.

Parameter	Data type	Type	Description	Default
user	String	required	Defines a DN used during the initial LDAP bind request. The {{root}} variable will be automatically substituted by the root DIT (as provided in the parent auth_ldap or auth_active_dir directive) before attempting to bind to the LDAP server.	
password	String	required	The password used during the initial LDAP bind request. The password may be specified as plaintext or as an absolute path to the file storing the secret. If using a file path, Shiny Server will retain root privileges under the assumption that the file is only readable by the root user.	

Applies to: [auth_ldap](#), [auth_active_dir](#)

Inheritable: Yes

user_bind_template

Define the template by which the username provided by the user should be converted to the username used to bind to the LDAP server. We currently only support storing all users in a single node of the directory.

Parameter	Data type	Type	Description	Default
template	String	required	The template used in defining the DN of users in the LDAP database when trying to bind. The following variables will be automatically substituted before attempting to bind to the LDAP server: {{username}} will be replaced by the username provided on the login page; {{root}} will be replaced by the root DIT (as provided in the parent auth_ldap directive). If this setting is not provided, defaults will be used. For auth_active_dir systems, the the default is '{{username}}@example.org' (where 'example.org' is the domain inferred/provided in auth_active_dir). For other auth_ldap systems, the default will be 'uid={{username}},ou=People,{{root}}'.	

Applies to: [auth_ldap](#), [auth_active_dir](#)
Inheritable: Yes

trusted_ca

By default, many commonly accepted Certificate Authorities (CAs) will be trusted when establishing secure ldap (ldaps://) connections. This parameter instructs Shiny Server to instead trust this CA certificate. Organizations that are using an internal CA to sign their SSL certificates or are using a self-signed SSL certificate on their LDAP server will need to provide the associated CA certificate in order for Shiny Server to trust the ldaps:// connection.

Parameter	Data type	Type	Description	Default
certificate	String	multiple	The path to the CA certificate to be trusted in PEM format.	

Applies to: [auth_ldap](#), [auth_active_dir](#)
Inheritable: Yes

check_ssl_ca

Define whether or not to check the Certificat Authority (CA) of the SSL certificate returned from the LDAP server. This must be true if using a self-signed SSL certificate. If not provided, this is set to 'true'. This should only be set to false if you know that your LDAP server uses an untrusted SSL certificate.

Parameter	Data type	Type	Description	Default
value	Boolean	required	If true, will check the certificate against the default list of trusted CAs (or the list provided in the trusted_ca setting). If false, the returned SSL certificate will not be checked.	

Applies to: [auth_ldap](#), [auth_active_dir](#)
Inheritable: Yes

user_filter

The filter to be used when querying the LDAP database to find the LDAP user object which contains the username. In Active Directory, the CN of the users are often not the username. In such instances, it is necessary to map the entered username to the user's DN. This will be done by querying in [user_search_base](#) for objects using this filter. If this field is left empty, it will be assumed that the name the user enters is the same as the associated LDAP user's CN, and can be used in the user's DN as instructed by [user_bind_template](#). The following variables are available: {{username}} will be replaced by the username provided on the login page; {{root}} will be replaced by the root DIT (as provided in the parent [auth_ldap](#) directive); {{userBind}} will be replaced by the username after applying the user_bind_template. The default for this setting in auth_ldap is to leave this field empty; the default in [auth_active_dir](#) is 'userPrincipalName={{userBind}}'

Parameter	Data type	Type	Description	Default
filter	String	required	The filter to use to find the LDAP user object referenced by the given username (as entered by the user on the login page).	

Applies to: [auth_ldap](#), [auth_active_dir](#)
Inheritable: Yes

user_search_base

Set the subtree which will be the root of user searches. If this setting is omitted, the defaults of 'cn=Users' or 'ou=People' will be used for systems using [auth_active_dir](#) and [auth_ldap](#), respectively. The given string will be followed by a comma and the root DIT to form the base of the search for users. If empty, the unaltered root DIT will be used. This setting will only be used if [user_filter](#) is set; otherwise, no queries for users will be executed.

Parameter	Data type	Type	Description	Default
base	String	required	The part of the tree which stores users.	

Applies to: [auth_ldap](#), [auth_active_dir](#)
Inheritable: Yes

group_searchn_base

Set the subtree which will be the root of group searches. If this setting is omitted, the defaults of 'cn=Users' or 'ou=Groups' will be used for systems using [auth_active_dir](#) and [auth_ldap](#), respectively. The given string will be followed by a comma and the root DIT to form the base of the search for groups. If empty, the unaltered root DIT will be used.

Parameter	Data type	Type	Description	Default
base	String	required	The part of the tree which stores groups.	

Applies to: [auth_ldap](#), [auth_active_dir](#)

Inheritable: Yes

group_filter

The LDAP query to use in determining group membership. The query should return all groups of which the given user is a member. Three variables are available for this query: '{{username}}' will be replaced with the username entered at the login screen; '{{userDN}}' will be replaced with the user's full DN (using the [user_bind_template](#) to compute this); '{{root}}' will be replaced by the root DIT (as provided in the parent [auth_ldap](#) directive). For [auth_active_dir](#), the default is 'member:1.2.840.1.13556.1.4.1941:={{userDN}}'. For [auth_ldap](#), the default is 'uniqueMember={{username}}'.

Parameter	Data type	Type	Description	Default
filter	String	required	The filter to use to determine group membership.	

Applies to: [auth_ldap](#), [auth_active_dir](#)

Inheritable: Yes

group_name_attribute

The attribute of a group which contains the group name.

Parameter	Data type	Type	Description	Default
attribute	String	optional	The name of the attribute. Default is 'cn' for both auth_active_dir and auth_ldap	cn

Applies to: [auth_ldap](#), [auth_active_dir](#)

Inheritable: Yes

application

DEPRECATED. This setting is deprecated and no longer enforced in Shiny Server. It is not permitted in configuration files.

Parameter	Data type	Type	Description	Default
empty	String	required		

Applies to: [location](#)

Inheritable: Yes

Child directives: [required_group](#), [required_user](#), [google_analytics_id](#), [app_init_timeout](#), [app_idle_timeout](#), [app_session_timeout](#), [simple_scheduler](#), [utilization_scheduler](#), [whitelist_headers](#)

template_dir

A directory containing custom templates to be used when generating pages in Shiny Server.

Parameter	Data type	Type	Description	Default
dir	String	required	The directory containing HTML templates.	

Applies to: Top-level, [server](#), [location](#), [admin](#)

Inheritable: Yes

pam_sessions_profile

When spawning a Shiny process for a user, the PAM profile specified here will be used. The default is to use the 'su' profile. You

can customize this behavior by creating a new profile in `/etc/pam.d` and referencing it here.

Parameter	Data type	Type	Description	Default
profile	String	required	The name of the profile in <code>/etc/pam.d</code> you want to use.	

Applies to: Top-level, [server](#), [location](#)
Inheritable: Yes

exec_supervisor

The program supervisor under which Shiny processes should be run. This can be a command that modifies Shiny's environment or resources. The default is to have no supervisor (i.e. the empty string `"`).

Parameter	Data type	Type	Description	Default
command	String	required	The command which prefixes the call to start the R process which will run Shiny.	

Applies to: Top-level, [server](#), [location](#)
Inheritable: Yes

r_path

The path to the R executable.

Parameter	Data type	Type	Description	Default
path	String	required	The absolute path on disk to the R executable you wish to use.	

Applies to: Top-level, [server](#), [location](#)
Inheritable: Yes

disable_websockets

Disable WebSockets on connections to the server. Some networks will not reliably support WebSockets, so this setting can be used to force Shiny Server to fall back to another protocol to communicate with the server. This is equivalent to adding `'websocket'` to `disabled_protocols`

Parameter	Data type	Type	Description	Default
val	Boolean	optional	Whether or not WebSockets should be disabled.	true

Applies to: Top-level
Inheritable: Yes

rrd_disabled

If present, RRD logging will be disabled. Note that having RRD is a prerequisite for enabling the [admin](#) interface. While this option is present, no data will be logged to be viewed in the Admin interface in the future, either.

This directive has no parameters.

Applies to: Top-level
Inheritable: Yes

disable_protocols

Disable some of the SockJS protocols used to establish a connection between your users and your server. Some network configurations cause problems with particular protocols; this option allows you to disable those.

Parameter	Data type	Type	Description	Default
names	String	multiple	The protocol(s) to disable. Available protocols are: <code>'websocket'</code> , <code>'xdr-streaming'</code> , <code>'xhr-streaming'</code> , <code>'iframe-eventsouce'</code> , <code>'iframe-htmlfile'</code> , <code>'xdr-polling'</code> , <code>'xhr-polling'</code> , <code>'iframe-xhr-polling'</code> , <code>'jsonp-polling'</code>	

Applies to: Top-level
Inheritable: Yes

graphite_enabled

Enable logging of metrics to a Graphite server.

Parameter	Data type	Type	Description	Default
host	String	optional	The address of the Graphite server. Defaults to the local server.	127.0.0.1
port	Integer	optional	The port on the Graphite server to which we should send metrics.	2003

Applies to: Top-level
Inheritable: Yes

7.3 Migrating & Updating

7.3.1 Upgrading from Shiny Server 0.3.x

Shiny Server 0.4 and later include many major changes to the architecture of Shiny Server. There are a few things to be aware of when upgrading from Shiny Server 0.3.x.

First, Shiny Server is now distributed via deb and rpm installers, rather than using npm. Thus, you'll first want to uninstall the old version of Shiny Server which was installed from npm.

```
$ sudo npm uninstall -g shiny-server
```

Shiny Server no longer requires node.js to be installed on the system either. If you have no other need for node.js, you can uninstall it at this time, as well.

Running the latest Shiny Server installer will replace the scripts at /etc/init/shiny-server.conf or /etc/init.d/shiny-server. If you have made any modifications to these startup scripts that you wish to keep, please save a copy of these files before running the installer.

The default directories for logging and Shiny application hosting have also moved in this release. The default logging directory used to be /var/shiny-server/log, it is now /var/log/shiny-server. The default hosting directory was previously /var/shiny-server/www and is now /srv/shiny-server/. If you have not explicitly overridden these directories in the /etc/shiny-server/shiny-server.conf file, the new location will be created and used in the latest Shiny Server.

Finally, Shiny Server 0.4 and later require the Shiny package to be at least version 0.7.0. You can check the system-wide version of Shiny you have installed using the following command on the command line.

```
$ sudo su - -c "R -e \"packageVersion('shiny')\""
```

(Running this command using the sudo su - -c preface will allow you to see the system-wide version of Shiny. Individual users could have installed more recent versions of Shiny in their own R libraries.) If the installed version of Shiny predates 0.7.0, you should follow the instructions in [Install Shiny](#) to update to the most recent version.

At this point, you can proceed with the [Installation](#) instructions associated with your Operating System.

7.3.2 Upgrading from Shiny Server 0.4.x

There are problems on some systems with updating from 0.4.x versions to a more recent version. Thus, any user updating from the 0.4.x line should uninstall the package before installing the new one. In all versions 0.5.x and later, updating can simply be done by installing the newer version over the old one without explicitly removing the old package.

7.4 Frequently Asked Questions

7.4.1 What are the hardware requirements for running Shiny Server?

The performance footprint of a Shiny application is almost entirely dependent upon the Shiny application code. There are two factors to consider when selecting the hardware platform for Shiny Server.

RAM

The memory requirements of a Shiny application depend heavily on the data loaded when running the application. Often users find that a Shiny R process requires a minimum of 50MB of RAM -- beyond this, the amount of memory consumed by an application is determined by the data loaded or generated by that application. The [Scoping Section](#) of the Shiny Tutorial describes in detail how to take advantage of Shiny's scoping rules to share data across multiple Shiny sessions. This enables application developers to **load only one copy of data into memory but still share this data with multiple shiny sessions.**

CPU

R is fundamentally a **single-threaded** application. Unless parallel packages and tools are specifically selected when designing a Shiny application, the average R process (and the associated Shiny application) will be run serially on a single processing core. Thus, the typical Shiny application may saturate the processing core to which it's assigned, but will be unable to leverage other cores on the server which may be idle at that time.

7.4.2 Does Shiny Server work on Amazon Elastic Compute Cloud (EC2)?

Absolutely. Many customers deploy Shiny Server and Shiny Server Professional on the Amazon Web Services (AWS) framework.

7.4.3 Can I Deploy Shiny Server on an Isolated Network?

Of course. Shiny Server does not require a connection to the Internet in order to work properly, so you are free to deploy it in whatever sort of network configuration you would like. **Offline** activation is also available for Shiny Server Professional customers.

7.4.4 How Do I Translate My 'application' Settings to Nested 'location's?

Shiny Server 0.4.2 introduced a new model for managing applications at a more granular level by supporting nested location directives. Any existing application directives in your configuration file will need to be migrated to this nested location format in order to be supported by any version of Shiny Server after 0.4.2.

application directives were addressed by their full path on disk. Nested locations, on the other hand, are indexed by their relative path within their parent location. For instance, the following exemplifies how you could redirect a particular application in the old model:

```
location / {
  site_dir /srv/shiny-server;

  # Define rules to apply to one application, stored at '/srv/shiny-server/app1'
  application /srv/shiny-server/app1 {
    redirect "http://rstudio.com" 302 true;
  }
}
```

In the new "nested locations" model, the equivalent would be:

```
location / {
  site_dir /srv/shiny-server;

  # Define rules to apply to one application, stored at '/srv/shiny-server/app1'
  location /app1 {
    redirect "http://rstudio.com" 302 true;
  }
}
```

This structure is much more powerful than simple application settings. The following configuration defines a parent location (depts) with some settings, then overrides those settings for a particular directory (finance) within that location.

```
server {
  ...
  # Define the '/depts' location
  location /depts {
    # Host a directory of applications
    site_dir /srv/departmentApps;
  }
}
```



```

# Provide a default/global GAID
google_analytics_id "UA-12345-1";

# Define the '/finance' location.
# Corresponds to the URL '/depts/finance', and the filesystem path
# '/srv/departmentApps/finance' as defined by the parent location.
location /finance {
    # Provide a custom GAID for only this sub-location
    google_analytics_id "UA-54321-9";

    location /fall2014/app1 {
        redirect "http://rstudio.com" 302 true;
    }
}
...

```

In this case, all applications stored in the `depts/finance/` directory would have a different Google Analytics ID than the rest of the applications on this server. Additionally, the application at `depts/finance/fall2014/app1` would be redirected to `http://rstudio.com`.

7.4.5 When I open a lot of Shiny Docs at once in Firefox, why do some not open?

Firefox limits the persistent connections that the browser will have open to a single server. For most use cases, this won't be a problem. But a user who opens multiple Shiny Docs with many embedded Shiny components all at once may start hitting this limit.

To resolve this, increase the [network.http.max-persistent-connections-per-server](#) setting (we tested by setting it to 50). Once that limit is increased, you should be able to open many complex Shiny Docs simultaneously from Firefox without issue.

7.4.6 Can I Bulk-Import in Flat-File Authentication? - Pro Only!

Some users may have an existing table of usernames and passwords they wish to import into a flat-file authentication system using the `sspasswd` tool. The shell script below can be used as a model for scripting such a solution, but be aware of the security concern mentioned below.

```

#!/bin/bash

PWFIL="passwdfile"
USERFILE="users.txt"

if [ ! -f "$PWFIL" ]; then
    touch $PWFIL
fi

while read p; do
    parts=( $p )
    echo "Adding ${parts[0]}..."
    echo "${parts[1]}" | /opt/shiny-server/bin/sspasswd \
        $PWFIL "${parts[0]}"
done < $USERFILE
echo "Done."

```

The script above will extract the usernames and passwords from the file specified in the `USERFILE` variable, and will use the file name provided in the `PWFIL` variable as the file in which the usernames and passwords should be stored. This script expects a user file which is stored in a tab-delimited format in which the username is in the first column, and the password is in the second.

Most secure password-generating tools recommend that the password be typed in an interactive console in which the text can be hidden, rather than passing it in via a command-line argument. There are (at least) two reasons for this:

1. Such commands (including the password, if it were provided as an argument to the command) are stored in the history file for your shell. Thus, a user with access to your shell history may be able to retrieve the password.
2. In a model like the one used above, the passwords won't be stored in your shell's history, but they will be available in the system's process table. The `echo` command above (and the password associated with it) will likely be accessible by any

other user on the system while that command is executing.

Be sure to fully understand the security implications of the above script before using it, especially if using it on an unsecured or multi-user server.

7.4.7 How Can I Set The trusted_ca For My LDAP Server Over SSL? - Pro Only!

The easiest way to confirm an SSL connection is to use the openssl tool to try to connect to your LDAP server. If you don't already have the SSL certificates for your server, you can even download them using this tool. If you run

```
openssl s_client -connect <LDAP server address>:<port> -showcerts
```

(the default LDAPS port is 636, by the way), you should get a flurry of output.

If you look through this output (in particular at the last few lines) you should see a "result"; if there's a problem it may say something like Verify return code: 19 (self signed certificate in certificate chain) which reveals that there's an issue with trusting the SSL connection between you and your LDAPS server. In the case of an error like the one above, you likely just need to instruct your client to trust a particular Certificate Authority (CA) that the openssl tool doesn't trust by default. Once you retrieve the CA certificate for your organization (which should also be the last certificate returned by the command above if you're actually connected to the right server), you can tell openssl to trust that CA by using a command in the format of

```
openssl s_client -connect <server-address>:<port> -CAfile <file.pem>
```

Assuming that the certificate matches the CA you provide (and everything is in the right format), you should get a line of output from openssl that says Verify return code: 0 (ok). Once you see that, you know you have your CA certificate in the right format.

There is one important check that the openssl tools doesn't perform that you should confirm before trying to use the certificate in Shiny Server Pro. You'll need to confirm that the hostname you're using matches the SSL certificate. You can do that manually or using curl by running `curl --cacert <file.pem> ldaps://<server-address>:<port>/`. If you see some LDAP output (perhaps starting with DN:) and no errors, then things are working properly and you have the right hostname.

Once you have the CA certificate working in the above tests, then you're ready to apply it to Shiny Server Pro. The CA certificate should be in PEM format and *only include one certificate per file*. You can add these file references using the trusted_ca setting in your Shiny Server Pro configuration as follows:

```
auth_ldap ... {  
  trusted_ca /etc/ssl/ca1.pem /etc/ssl/ca2.pem;  
}
```

This would include two CA certificates that Shiny Server Pro should trust. Shiny Server Pro should now be able to successfully connect to your LDAPS server when you attempt to authenticate users.